

- [301] V. Ganti, J. E. Gehrke, and R. Ramakrishnan. Cactus—clustering categorical data using summaries. In *Proc. ACM Intl. Conf. on Knowledge Discovery in Databases*, 1999.
- [302] V. Ganti, R. Ramakrishnan, T. E. Gehrke, A. Powell, and J. French. Clustering large datasets in arbitrary metric spaces. In *Proc. IEEE Intl. Conf. Data Engineering*, 1999.
- [303] H. Garcia-Molina and D. Barbara. How to assign votes in a distributed system. *Journal of the ACM*, 32(4), 1985.
- [304] H. Garcia-Molina, R. Lipton, and J. Valdes. A massive memory system machine. *IEEE Transactions on Computers*, C33(4):391–399, 1984.
- [305] H. Garcia-Molina, J. Ullman, and J. Widom. *Database Systems: The Complete Book*. Prentice Hall, 2001.
- [306] H. Garcia-Molina and G. Wiederhold. Read-only transactions in a distributed database. *ACM Transactions on Database Systems*, 7(2):209–234, 1982.
- [307] E. Garfield. Citation analysis as a tool in journal evaluation. *Science*, 178(4060):471–479, 1972.
- [308] A. Garg and C. Gottlieb. Order preserving key transformations. *ACM Transactions on Database Systems*, 11(2):213–234, 1986.
- [309] J. E. Gehrke, V. Ganti, R. Ramakrishnan, and W.-Y. Loh. Boat: Optimistic decision tree construction. In *Proc. ACM SIGMOD Conf. on Management of Data*, 1999.
- [310] J. E. Gehrke, F. Korn, and D. Srivastava. On computing correlated aggregates over continual data streams. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 2001.
- [311] J. E. Gehrke, R. Ramakrishnan, and V. Ganti. Rainforest: A framework for fast decision tree construction of large datasets. In *Proc. Intl. Conf. on Very Large Databases*, 1998.
- [312] S. P. Ghosh. *Data Base Organization for Data Management (2nd ed.)*. Academic Press, 1986.
- [313] P. B. Gibbons, Y. Matias, and V. Poosala. Fast incremental maintenance of approximate histograms. In *Proc. of the Conf. on Very Large Databases*, 1997.
- [314] P. B. Gibbons and Y. Matias. New sampling-based summary statistics for improving approximate query answers. In *Proc. ACM SIGMOD Conf. on the Management of Data*, pages 331–342. ACM Press, 1998.
- [315] D. Gibson, J. M. Kleinberg, and P. Raghavan. Clustering categorical data: An approach based on dynamical systems. In *Proc. Intl. Conf. Very Large Data Bases*, 1998.
- [316] D. Gibson, J. M. Kleinberg, and P. Raghavan. Inferring web communities from link topology. In *Proc. ACM Conf. on Hypertext*, 1998.
- [317] G. A. Gibson. *Redundant Disk Arrays: Reliable: Parallel Secondary Storage*. An ACL Distinguished Dissertation 1991. MIT Press, 1992.
- [318] D. Gifford. Weighted voting for replicated data. In *ACM Symp. on Operating Systems Principles*, 1979.
- [319] A. C. Gilbert, Y. Kotidis, S. Muthukrishnan, and M. J. Strauss. Surfing wavelets on streams: One-pass summaries for approximate aggregate queries. In *Proc. of the Conf. on Very Large Databases*, 2001.
- [320] C. F. Goldfarb and P. Prescod. *The XML Handbook*. PrenticeHall, 1998.
- [321] R. Goldman and L. Vidor. DataGuides: enabling query formulation and optimization in semistructured databases. In *Proc. Intl. Conf. on Very Large Data Bases*, pages 436–445, 1997.

- [322] J. Goldstein, R. Ramakrishnan, U. Shaft, and J.-B. Yu. Processing queries by linear constraints. In *Proc. ACM Symposium on Principles of Database Systems*, 1997.
- [323] G. Graefe. Encapsulation of parallelism in the Volcano query processing system. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1990.
- [324] G. Graefe. Query evaluation techniques for large databases. *ACM Computing Surveys*, 25(2), 1993.
- [325] G. Graefe, R. Bunker, and S. Cooper. Hash joins and hash tables in Microsoft SQL Server. In *Proc. Intl. Conf. on Very Large Databases*, 1998.
- [326] G. Graefe and D. DeWitt. The Exodus optimizer generator. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1987.
- [327] G. Graefe and K. Ward. Dynamic query optimization plans. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1989.
- [328] M. Graham, A. Mendelzon, and M. Vardi. Notions of dependency satisfaction. *Journal of the ACM*, 33(1):105-129, 1986.
- [329] G. Grahne. *The Problem of Incomplete Information in Relational Databases*. Springer-Verlag, 1991.
- [330] L. Gravano, H. Garcia-Violina, and A. Tornasice. Gloss: text-source discovery over the internet. *ACM Transactions on Database Systems*, 24(2), 1999.
- [331] J. Gray. Notes on database operating systems. In *Operating Systems: An Advanced Course*. eds. Bayer, Graham, and Seegmuller, Springer-Verlag, 1978.
- [332] J. Gray. The transaction concept: Virtues and limitations. In *Proc. Intl. Conf. on Very Large Databases*, 1981.
- [333] J. Gray. Transparency in its place—the case against transparent access to geographically distributed data. *Tandem Computers*, TR-8.9-1, 1989.
- [334] J. Gray. *The Benchmark Handbook: for Database and Transaction Processing Systems*. Morgan Kaufmann, 1991.
- [335] J. Gray, A. Bosworth, A. Layman, and H. Pirahesh. Data cube: A relational aggregation operator generalizing group-by, cross-tab and sub-totals. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1996.
- [336] J. Gray, R. Lorie, G. Putzolu, and I. Traiger. Granularity of locks and degrees of consistency in a shared database. In *Proc. of IFIP Working Conf. on Modelling of Data Base Management Systems*, 1977.
- [337] J. Gray, P. McJones, M. Blasgen, B. Lindsay, R. Lorie, G. Putzolu, T. Price, and I. Traiger. The recovery manager of the System R database manager. *ACM Computing Surveys*, 13(2):223-242, 1981.
- [338] J. Gray and A. Reuter. *Transaction Processing: Concepts and Techniques*. Morgan Kaufmann, 1992.
- [339] P. Gray. *Logic, Algebra, and Databases*. John Wiley, 1984.
- [340] M. Greenwald and S. Khanna. Space-efficient online computation of quantile summaries. In *Proc. ACM SIGMOD Conf. on Management of Data*, 2001.
- [341] P. Griffiths and B. Wade. An authorization mechanism for a relational database system. *ACM Transactions on Database Systems*, 1(3):242-255, 1976.
- [342] G. Grinstein. Visualization and data mining. In *Intl. Conf. on Knowledge Discovery in Databases*, 1996.

- [343] S. Guha, N. Mishra, R. Motwani, and L. O'CaHaghan. Clustering data streams. In *Proc. of the Annual Symp. on Foundations of Computer Science*, 2000.
- [344] S. Guha, R. Rastogi, and K. Shill. Cure: an efficient clustering algorithm for large databases. In *Proc. ACM SIGMOD Conf. on Management of Data*, 1998.
- [345] S. Guha, N. Koudas, and K. Shiri. Data streams and histograms. In *Proc. of the ACM Symp. on Theory of Computing*, 2001.
- [346] D. Gunopulos, H. Mannila, R. Khardon, and H. Toivonen. Data mining, hypergraph transversals, and machine learning. In *Proc. ACM Symposium on Principles of Database Systems*, pages 209-216, 1997.
- [347] D. Gunopulos, H. Mannila, and S. Saluja. Discovering all most specific sentences by randomized algorithms. In *Proc. of the Intl. Conf. on Database Theory*, volume 1186 of *Lecture Notes in Computer Science*, pages 215-229. Springer, 1997.
- [348] A. Gupta and I. Munlick. *Materialized Views: Techniques, Implementations, and Applications* MIT Press, 1999.
- [349] A. Gupta, I. Munlick, and V. Subrahmanian. Maintaining views incrementally. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1993.
- [350] A. Guttman. R-trees: a dynamic index structure for spatial searching. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1984.
- [351] L. Haas, W. Chang, G. Lohman, J. McPherson, P. Wilms, G. Lapis, B. Lindsay, H. Pirahesh, M. Carey, and E. Shekita. Starburst mid-flight: As the dust clears. *IEEE Transactions on Knowledge and Data Engineering*, 2(1), 1990.
- [352] P. Haas, J. Naughton, S. Seshadri, and L. Stokes. Sampling-based estimation of the number of distinct values of an attribute. In *Proc. Intl. Conf. on Very Large Databases*, 1995.
- [353] P. Haas and A. Swami. Sampling-based selectivity estimation for joins using augmented frequent value statistics. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1995.
- [354] P. J. Haas and J. M. Hellerstein. Ripple joins for online aggregation. In *Proc. ACM SIGMOD Conf. on the Management of Data*, pages 287-298. ACM Press, 1999.
- [355] T. Haerder and A. Reuter. Principles of transaction oriented database recovery-a taxonomy. *ACM Computing Surveys*, 15(4), 1982.
- [356] U. Halici and A. Dogac. Concurrency control in distributed databases through time intervals and short-term locks. *IEEE Transactions on Software Engineering*, 15(8):994-1003, 1989.
- [357] M. Hall. *COTE Web Programming: IfTML, Java, CGI, & Javascript*. Prentice-Hall, 1997.
- [358] P. Hall. Optimization of a simple expression in a relational database system. *IBM Journal of Research and Development*, 20(3):244-257, 1976.
- [359] G. Hamilton, R. G. Cattell, and M. Fisher. *JDBC Database Access With Java: A Tutorial and Annotated Reference*. Java Series. Addison-Wesley, 1997.
- [360] M. Harnrner and D. McLeod. Semantic integrity in a relational database system. In *Proc. Intl. Conf. on Very Large Databases*, 1975.
- [361] J. Han and Y. Fu. Discovery of multiple-level association rules from large databases. In *Proc. Intl. Conf. on Very Large Databases*, 1995.
- [362] D. Hanel. *Construction and Assessment of Classification Rules*. John Wiley & sons, Chichester, England, 1997.

- [363] J. Han and M. Kamber. *Data Mining: Concepts and Techniques*. Morgan Kaufmann Publishers, 2000.
- [364] J. Han, J. Pei, and Y. Yin. Mining frequent patterns without candidate generation. In *Proc. ACM SIGMOD Int'l. Conf. on Management of Data*, pages 1–12, 2000.
- [365] E. Hanson. A performance analysis of view materialization strategies. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1987.
- [366] E. Hanson. Rule condition testing and action execution in Ariel. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1992.
- [367] V. Harinarayan, A. Rajaraman, and J. Ullman. Implementing data cubes efficiently. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1996.
- [368] J. Haritsa, M. Carey, and M. Livny. On being optimistic about real-time constraints. In *ACM Symp. on Principles of Database Systems*, 1990.
- [369] J. Harrison and S. Dietrich. Maintenance of materialized views in deductive databases: An update propagation approach. In *Proc. Workshop on Deductive Databases*, 1992.
- [370] T. Hastie, R. Tibshirani, and J. H. Friedman. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer Verlag, 2001.
- [371] D. Heckerman. Bayesian networks for knowledge discovery. In *Advances in Knowledge Discovery and Data Mining*, eds. U.M. Fayyad, G. Piatetsky-Shapiro, P. Smyth, and R. Uthurusamy, MIT Press, 1996.
- [372] D. Heckerman, H. Mannila, D. Pregibon, and R. Uthurusamy, editors. *Proc. Intl. Conf. on Knowledge Discovery and Data Mining*. AAAI Press, 1997.
- [373] J. Hellerstein. Optimization and execution techniques for queries with expensive methods. *Ph.D. thesis, University of Wisconsin-Madison*, 1995.
- [374] J. Hellerstein, P. Haas, and H. Wang. Online aggregation. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1997.
- [375] J. Hellerstein, E. Koutsoupias, and C. Papadimitriou. On the analysis of indexing schemes. In *Proceedings of the ACM Symposium on Principles of Database Systems*, pages 249–256. ACM Press, 1997.
- [376] J. Hellerstein, J. Naughton, and A. Pfeffer. Generalized search trees for database systems. In *Proc. Int'l. Conf. on Very Large Databases*, 1995.
- [377] J. M. Hellerstein, E. Koutsoupias, and C. H. Papadimitriou. On the analysis of indexing schemes. In *Proc. ACM Symposium on Principles of Database Systems*, pages 249–256, 1997.
- [378] C. H. Hsiao. Online association rule mining. In *Proc. ACM SIGMOD Conf. on the Management of Data*, pages 145–156, 1999.
- [379] R. Himmelfrucht, G. Lausen, B. Ludascher, and C. Schlepphorst. On a declarative semantics for Web queries. *Lecture Notes in Computer Science*, 1341:386–398, 1997.
- [380] C.-Y. Ho, R. Agrawal, N. Megiddo, and R. Srikant. Range queries in OLAP data cubes. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1997.
- [381] S. Holzner. *XML Complete*. McGraw-Hill, 1998.
- [382] D. Hong, T. Johnson, and U. Chakravarthy. Real-time transaction scheduling: A cost conscious approach. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1993.
- [383] W. Hong and M. Stonebraker. Optimization of parallel query execution plans in XPRS. In *Proc. Intl. Conf. on Parallel and Distributed Information Systems*, 1991.

- [384] W.-C. HOLL and G. Ozsoyoglu. Statistical estimators for aggregate relational algebra queries. *ACM Transactions on Database Systems*, 16(4), 1991.
- [385] H. Hsiao and D. DeWitt. A performance study of three high availability data replication strategies. In *Proc. Intl. Conf. on Parallel and Distributed Info'rmation Systems*, 1991.
- [386] J. Huang, J. Stankovic, K. RamallrithaIll, and D. Towsley. Experimental evaluation of real-tirne optiInistic concurrency control SChellles. In *Proc. Intl. Conf. on Very Large Databases*, 1991.
- [387] Y. Huang, A. Sistla, and O. vVolfson. Data replication for rrObile cOlnputers. In *Proc. ACM SIGMOD Conj. on the Management of Data*, 1994.
- [388] Y. Huang and O. Wolfson. A cOlnpetitive dynarnic data replication algorithm. In *Proc. IEEE CS IEEE Inti. Conf. on Data Engineering*, 1993.
- [389] R. Hull. Managing semantic heterogeneity in databases: A theoretical perspective. In *ACM Symp. on Principles of Database Syste'ns*, 1997.
- [390] R. Hull and R. King. Semantic database modeling: Survey, applications, and research issues. *ACM Cornputing Surveys*, 19(19):201-260, 1987.
- [391] R. Hull and J. Suo Algebraic and calculus query languages for recursively typed complex objects. *Journal of Computer and System Sciences*, 47(1):121-156, 1993.
- [392] R. Hull and M. Yoshikawa. ILOG: Declarative creation and rnanipulation of object-identifiers. In *Proc. Inti. Conf. on Very Large Databases*, 1990.
- [393] G. Hulten, L. Spencer, and P. Domingos. Mining tillle-changing data streamns. In *Proc. ACM SIGKDD Intl. Conference on Knowledge Discovery and Data M/Jining*, pages 97-106. AAAI Press, 2001.
- [394] J. Hunter. *Java Servlet Programming*. O'Reilly Associates, Inc., 1998.
- [395] T. Imielinski and H. Korth (eds.). *Mobile Computing*. Kluwer Academic, 1996.
- [396] T. Imielinski and W. Lipski. Incomplete information in relational databases. *Journal of the ACM*, 31(4):761-791, 1984.
- [397] T. Imielinski and H. Mannila. A database perspective on knowledge discovery. *Communications of the ACM*, 38(11):58-64, 1996.
- [398] T. Imielinski, S. Viswanathan, and B. Badrinath. Energy efficient indexing on air. In *Proc. ACM SIGJv!OD Conf. on the Management of Data*, 1994.
- [399] Y. Ioannidis. Query optimization. In *Handbook of Comp'uteT Science*. ed. A.B. Tucker, CRC Press, 1996.
- [400] Y. Ioannidis and S. Christodoulakis. Optirnal histograms for liIlliting worst-case error propagation in the size of join results. *ACM Transactions on Database Systems*, 1993.
- [401] Y. Ioannidis and Y. Kang. Randornized algorithms for optimizing large join queries. In *Proc. ACM SIGMOD Conf. on the Jyfanagement of Data*, 1990.
- [402] Y. Ioannidis and Y. Kang. Ileft-deep vs. bushy trees: An analysis of strategy spaces and its irnplications for query optirnization. In *PROC. ACM SIGMOD Conf. on the Management of Data*, 1991.
- [403] Y. Ioannidis, R. Ng, K. Shirn, and T. Sellis. Parallletric query processing. In *Proc. Ina. Conf. on Very Large Databases*, 1992.
- [404] Y. Ioannidis and R. Rarnakrishnan. Containment of conjunctive queries: Beyond relations as sets. *ACM TransactioTls on Database Sy8terns*, 20(3):288-324, 1995.
- [405] Y. E. Ioannidis. Universality of serial histograms. In *Proc. Intl. Conf. on Ve'ry Large Database8*, 199:3.

- [406] H. Jagadish, D. Lieuwen, R. Rastogi, A. Silberschatz, and S. Sudarshan. Dali: A high performance in-memory storage manager. In *Proc. Intl. Conf. on Very Large Databases*, 1994.
- [407] A. K. Jain and R. C. Dubes. *Algorithms for Clustering Data*. PrenticeHall, 1988.
- [408] S. Jajodia and D. Mutchler. Dynamic voting algorithms for maintaining the consistency of a replicated database. *ACM Transactions on Database Systems*, 15(2):230–280, 1990.
- (409) S. Jajodia and R. Sandhu. Polyinstantiation integrity in multilevel relations. In *Proc. IEEE Symp. on Security and Privacy*, 1990.
- [410] M. Jarke and J. Koch. Query optimization in database systems. *ACM Computing Surveys*, 16(2):111–152, 1984.
- [411] K. S. Jones and P. Willett, editors. *Readings in Information Retrieval*. Multimedia Information and Systems. Morgan Kaufmann Publishers, 1997.
- [412] J. Lou and P. Fischer. The complexity of recognizing 3NF schemes. *Information Processing Letters*, 14(4):187–190, 1983.
- [413] N. Kabra and D. J. DeWitt. Efficient mid-query re-optimization of sub-optimal query execution plans. In *Proc. ACM SIGMOD Intl. Conf. on Management of Data*, 1998.
- [414] Y. Kambayashi, M. Yoshikawa, and S. Yajima. Query processing for distributed databases using generalized semi-joins. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1982.
- [415] P. Kanellakis. Elements of relational database theory. In *Handbook of Theoretical Computer Science*. ed. J. Van Leeuwen, Elsevier, 1991.
- [416] P. Kanellakis. Constraint programming and database languages: A tutorial. In *ACM Symp. on Principles of Database Systems*, 1995.
- [417] H. Kargupta and P. Chan, editors. *Advances in Distributed and Parallel Knowledge Discovery*. MIT Press, 2000.
- [418] L. Kaufman and P. Rousseeuw. *Finding Groups in Data: An Introduction to Cluster Analysis*. John Wiley and Sons, 1990.
- [419] R. Kaushik, P. Bohannon, J. F. Naughton, and H. F. Korth. Covering indexes for branching path expression queries. In *Proceedings of SIGMOD*, 2002.
- [420] D. Keim and H.-P. Kriegel. VisDB: a system for visualizing large databases. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1995.
- [421] D. Keim and H.-P. Kriegel. Visualization techniques for mining large databases: A comparison. *IEEE Transactions on Knowledge and Data Engineering*, 8(6):923–938, 1996.
- [422] A. Keller. Algorithms for translating view updates to database updates for views involving selections, projections, and joins. *ACM Symp. on Principles of Database Systems*, 1985.
- [423] W. Kent. *Data and Reality, Basic Assumptions in Data Processing Reconsidered*. North-Holland, 1978.
- [424] W. Kent, R. Ahmed, J. Albert, M. Ketabchi, and M.-C. Shan. Object identification in multi-database systems. In *IFIP Intl. Conf. on Data Semantics*, 1992.
- [425] L. Kerschberg, A. Klug, and D. Tsichritzis. A taxonomy of data models. In *Systems for Large Data Bases*. eds. P.C. Lockemann and E.J. Neuhold, North-Holland, 1977.
- [426] W. Kiessling. On semantic reefs and efficient processing of correlation queries with aggregates. In *Proc. Intl. Conf. on Very Large Databases*, 1985.

- [427] M. Kifer, V. Kim, and Y. Sagiv. Querying object-oriented databases. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1992.
- [428] M. Kifer, G. Lausen, and J. Wu. Logical foundations of object-oriented and frame-based languages. *Journal of the ACM*, 42(4):741–843, 1995.
- [429] M. Kifer and E. Lozinskii. Sygraf: Implementing logic programs in a database style. *IEEE Transactions on Software Engineering*, 14(7):922–935, 1988.
- [430] W. Kim. On optimizing an SQL-like nested query. *ACM Transactions on Database Systems*, 7(3), 1982.
- [431] W. Kim. Object-oriented database systems: Promise, reality, and future. In *Proc. Intl. Conf. on Very Large Databases*, 1993.
- [432] W. Kim, J. Garza, N. Ballou, and D. Woelk. Architecture of the ORION next-generation database system. *IEEE Transactions on Knowledge and Data Engineering*, 2(1):109–124, 1990.
- [433] W. Kim and F. Lochovsky (eds.). *Object-Oriented Concepts, Databases, and Applications*. Addison-Wesley, 1989.
- [434] W. Kim, D. Reiner, and D. Batory (eds.). *Query Processing in Database Systems*. Springer Verlag, 1984.
- [435] W. Kim (ed.). *Modern Database Systems*. ACM Press and Addison-Wesley, 1995.
- [436] R. Kimball. *The Data Warehouse Toolkit*. John Wiley and Sons, 1996.
- [437] J. King. Quist: A system for semantic query optimization in relational databases. In *PTOC. Intl. Conf. on Very Large Databases*, 1981.
- [438] J. M. Kleinberg. Authoritative sources in a hyperlinked environment. In *Proc. ACM-SIAM Symp. on Discrete Algorithms*, 1998.
- [439] A. Klug. Equivalence of relational algebra and relational calculus query languages having aggregate functions. *Journal of the ACM*, 29(3):699–717, 1982.
- [440] A. Klug. On conjunctive queries containing inequalities. *Journal of the ACM*, 35(1):146–160, 1988.
- [441] E. Knapp. Deadlock detection in distributed databases. *ACM Computing Surveys*, 19(4):303–328, 1987.
- [442] D. Knuth. *The Art of Computer Programming, Vol. 3: Sorting and Searching*. Addison-Wesley, 1973.
- [443] G. Koch and K. Loney. *Oracle: The Complete Reference*. Oracle Press, Osborne-McGraw-Hill, 1995.
- [444] W. Kohler. A survey of techniques for synchronization and recovery in decentralized computer systems. *ACM Computing Surveys*, 13(2):149–184, 1981.
- [445] D. Konopnicki and O. Shmueli. W3QS: A system for WWW querying. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1997.
- [446] F. Korn, H. Jagadish, and C. Faloutsos. Efficiently supporting ad hoc queries in large datasets of time sequences. In *PTOC. ACM SIGMOD Conf. on Management of Data*, 1997.
- [447] M. Kornacker, C. Fvlohan, and J. Hellerstein. Concurrency and recovery in generalized search trees. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1997.
- [448] H. Korth, N. Soparkar, and A. Silberschatz. Triggered real-time databases with consistency constraints. In *Proc. Intl. Conf. on Very Large Databases*, 1990.

- [449] H. F. Korth. Deadlock freedom using edge locks. *ACM Transactions on Database Systems*, 7(4):632-652, 1982.
- [450] D. Kossman. The state of the art in distributed query processing. *ACM Computing Surveys*, 32(4):422-469, 2000.
- [451] Y. Kotidis and N. Roussopoulos. An alternative storage organization for ROLAP aggregate views based on cubetrees. In *Proc. ACM SIGMOD Intl. Conf. on Management of Data*, 1998.
- [452] N. Krishnakummar and A. Bernstein. High throughput escrow algorithms for replicated databases. In *Proc. Intl. Conf. on Very Large Databases*, 1992.
- [453] R. Krishnamurthy, H. Borat, and C. Zaniolo. Optimization of nonrecursive queries. In *Proc. Intl. Conf. on Very Large Databases*, 1986.
- [454] J. Kuhns. Logical aspects of question answering by computer. Technical report, Rand Corporation, RM-5428-Pr., 1967.
- [455] V. Kumar. *Performance of Concurrency Control Mechanisms in Centralized Database Systems*. PrenticeHall, 1996.
- [456] H. Kung and P. Lehman. Concurrent manipulation of binary search trees. *ACM Transactions on Database Systems*, 5(3):354-382, 1980.
- [457] H. Kung and J. Robinson. On optimistic methods for concurrency control. *PTOC. Intl. Conf. on Very Large Databases*, 1979.
- [458] D. Kuo. Model and verification of a data manager based on ARIES. In *Intl. Conf. on Database Theory*, 1992.
- [459] M. LaCroix and A. Pirotte. Domain oriented relational languages. In *PTOC. Intl. Conf. on Very Large Databases*, 1977.
- [460] M.-Y. Lai and W. Wilkinson. Distributed transaction management in Jasmin. In *Proc. Intl. Conf. on Very Large Databases*, 1984.
- [461] L. Lakshmanan, F. Sadri, and I. N. Subramanian. A declarative query language for querying and restructuring the web. In *Proc. Intl. Conf. on Research Issues in Data Engineering*, 1996.
- [462] L. V. S. Lakshmanan, Raymond T. Ng, J. Han, and A. Pang. Optimization of constrained frequent set queries with 2-variable constraints. In *Proc. ACM SIGMOD Intl. Conf. on Management of Data*, pages 157-168. ACM Press, 1999.
- [463] C. Larn, G. Landis, J. Orenstein, and D. Weinreb. The Objectstore database system. *Communications of the ACM*, 34(10), 1991.
- [464] L. Lamport. Time, clocks and the ordering of events in a distributed system. *Communications of the ACM*, 21(7):558-565, 1978.
- [465] B. Lampson and D. Lornet. A new presumed commit optimization for two phase commit. In *Proc. Intl. Conf. on Very Large Databases*, 1993.
- [466] B. Lampson and H. Sturgis. Crash recovery in a distributed data storage system. Technical report, Xerox PARC, 1976.
- [467] C. Landwehr. Formal models of computer security. *ACM Computing Surveys*, 13(3):247-278, 1981.
- [468] R. Langerak. View updates in relational databases with an independent scheme. *ACM Transactions on Database Systems*, 15(1):40-66, 1990.
- [469] P.-A. Larson. Linear hashing with overflow-handling by linear probing. *ACM Transactions on Database Systems*, 10(1):75-89, 1985.

- [470] P.-A. Larson. Linear hashing with separators—A dynamic hashing scheme achieving one-access retrieval. *ACM Transactions on Database Systems*, 13(3):366–388, 1988.
- [471] P.-A. Larson and G. Graefe. Memory Management During Run Generation in External Sorting. In *Proc. ACM SIGMOD Conf. on Management of Data*, 1998.
- [472] P. Lehman and S. Yao. Efficient locking for concurrent operations on B trees. *ACM Transactions on Database Systems*, 6(4):650–670, 1981.
- [473] T. Leung and R. Muntz. Temporal query processing and optimization in multiprocessor database machines. In *Proc. Intl. Conf. on Very Large Databases*, 1992.
- [474] M. Leventhal, D. Lewis, and M. Fuchs. *Designing Xylog Internet applications*. The Charles F. Goldfarb series on open information management. PrenticeHall, 1998.
- [475] P. Lewis, A. Bernstein, and M. Kifer. *Databases and Transaction Processing*. Addison Wesley, 2001.
- [476] E.-P. Lim and J. Srivastava. Query optimization and processing in federated database systems. In *Proc. Intl. Conf. on Intelligent Knowledge Management*, 1993.
- [477] B. Lindsay, J. McPherson, and H. Pirahesh. A data management extension architecture. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1987.
- [478] B. Lindsay, P. Selinger, C. Galtieri, J. Gray, R. Lorie, G. Putzolu, I. Traiger, and B. Wade. Notes on distributed databases. Technical report, RJ2571, San Jose, CA, 1979.
- [479] D.-I. Lin and Z. M. Kedem. Pincer search: A new algorithm for discovering the maximum frequent set. *Lecture Notes in Computer Science*, 1377:105–77, 1998.
- [480] V. Linnemann, K. Kuspert, P. Dada, P. Pistol, R. Erbe, A. Kenlper, N. Sudkamp, G. Walch, and M. Wallrath. Design and implementation of an extensible database management system supporting user defined data types and functions. In *Proc. Intl. Conf. on Very Large Databases*, 1988.
- [481] R. Lipton, J. Naughton, and D. Schneider. Practical selectivity estimation through adaptive sampling. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1990.
- [482] B. Liskov, A. Adya, M. Castro, M. Day, S. Ghemawat, R. Gruber, U. Maheshwari, A. Myers, and L. Shrira. Safe and efficient sharing of persistent objects in Thor. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1996.
- [483] W. Litwin. Linear Hashing: A new tool for file and table addressing. In *Proc. Intl. Conf. on Very Large Databases*, 1980.
- [484] W. Litwin. Trie Hashing. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1981.
- [485] W. Litwin and A. Abdellatif. Multidatabase interoperability. *IEEE Computer*, 12(19):10–18, 1986.
- [486] W. Litwin, L. Mark, and N. Roussopoulos. Interoperability of multiple autonomous databases. *ACM Computing Surveys*, 22(3), 1990.
- [487] W. Litwin, M.-A. Neirnat, and D. Schneider. LH^{*}—A scalable, distributed data structure. *ACM Transactions on Database Systems*, 21(4):480–525, 1996.
- [488] M. Liu, A. Sheth, and A. Singhal. An adaptive concurrency control strategy for distributed database system. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1984.
- [489] M. Livny, R. Rarnakrishnan, K. Beyer, G. Chen, D. Donjerkovic, S. Lawande, J. Myllyrnaki, and K. Wenger. DEVise: Integrated querying and visual exploration of large datasets. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1997.

- [490] G. Lohrnau. Granuar-like functional rules for representing query optiInization alternatives. In *ProG. ACM SIGMOD ConI on the Management of Data*, 1988.
- [491] D. Lomet and B. Salzberg. The hB-T ree: A rnultiattribute indexing lllethod with good guaranteed performance. *ACM Transactions on Database Systems*, 15(4), 1990.
- [492] D. Lomet and B. Salzberg. Access method concurrency with recovery. In *ProG. ACM SIGMOD Conf. on the Management of Data*, 1992.
- [493] R. Lorie. Physical integrity in a large segnlented database. *ACM Transactions on Database Systems*, 2(1):91-104, 1977.
- [494] R. Lorie and H. Young. A low COlllllllUnication sort algorithm for a parallel database machine. In *ProG. Intl. Conf. on Very Large Database.s*, 1989.
- [495] Y. Lou and Z. Ozsoyoglu. LLO: An object-oriented deductive language with methods and method inheritance. In *ProG. ACM SIGlvIOD Conf. on the Management of Data*, 1991.
- [496] H. Lu, B.-C. Ooi, and K.-L. Tan (eds.). *Query Processing in Parallel Relat'ional Database Systems*. IEEE Computer Society Press, 1994.
- [497] C. Lucchesi and S. Osborn. Candidate keys for relations. *J. Computer and System Sciences*, 17(2):270-279, 1978.
- [498] V. Lum. Multi-attribute retrieval with combined indexes. *Communications of the ACM*, 1(11):660-665, 1970.
- [499] T. Lunt, D. Denning, R. Schell, M. Heckman, and W. Shockley. The seaview security lllodel. *IEEE Transactions on Software Engineering*, 16(6):593---607, 1990.
- [500] L. Mackert and G. Lohrnau. R* optimizer validation and performance evaluation for local queries. Technical report, IBM RJ-4989, San Jose, CA, 1986.
- [501] D. Maier. *The Theory of Relational Databases*. Computer Science Press, 1983.
- [502] D. l\laier, A. Mendelzon, and Y. Sagiv. Testing irnplication of data dependencies. *ACM Transactions on Database Systems*, 4(4), 1979.
- [503] D. Maier and D. Warren. *Cornputing with Logic: Logic Programming with Prolog*. BenjaminjCurnnlings Publishers, 1988.
- [504] A. Makinouchi. A consideration on normal fonn of not-necessarily-nonnalized relation in the relational data rmodel. In *Proc. Intl. Conf. on Very Large Databases*, 1977.
- [505] U. l\lanber and R. Ladner. Concurrency control in a dynalnic search structure. *ACM Transactions on Database Systems*, 9(3):439--455, 1984.
- [506] G. l\lanku, S. Rajagopalan, and B. Lindsay. Handolll salnpling techniques for space efficient online COITlputation of order statistics of large datasets. In *Proc. ACM SIGJvfOD Conf. on Nfanagement of Data*, 1999.
- [507] H. Ivlannila. Methods and problems in data nlining. In *Intl. Conf. on Database Theory*, 1997.
- [508] H. Mannila and K.-J. Raiha. Design by Exarnple: An application of ArlTIstrong relations. *Journal of Computer and System Sciences*, 3:3(2):126--141, 1986.
- [509] H. rvIannila and K.-J. Raiha. *The Design of Relational Databases*. Addison-Wesley, 1992.
- [510] H. Mannila, H. Toivonen, and A. l. VerkarTlo. Discovering frequent episodes in sequences. In *proc. Intl. Conf. on Kno'wledge Discovery in Databases and Data Mining*, 1995.
- [511] H. Mannila, P. Sillyth, and D. J. Hanel. *Principles of Data Mining*. MIT Press, 20tH.

- [512] M. Mannino, P. Chu, and T. Sager. Statistical profile estimation in database SystCIUS. *ACM Computing Surveys*, 20(3):191–221, 1988.
- [513] V. Ilarkowitz. Representing processes in the extended entity-relationship model. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1990.
- [514] V. Markowitz. Safe referential integrity structures in relational databases. In *Proc. Intl. Conf. on Very Large Databases*, 1991.
- [515] Y. Matias, J. S. Vitter, and M. Vang. Dynamic maintenance of wavelet-based histograms. In *Proc. of the Conf. on Very Large Databases*, 2000.
- [516] D. McCarthy and U. Dayal. The architecture of an active data base management system. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1989.
- [517] W. McCune and L. Henschen. Maintaining state constraints in relational databases: A proof theoretic basis. *Journal of the ACM*, 36(1):46–68, 1989.
- [518] J. McHugh, S. Abiteboul, R. Goldman, D. Quass, and J. Widom. Lore: A database management system for semistructured data. *ACM SIGMOD Record*, 26(3):54–66, 1997.
- [519] S. Mehrotra, R. Rastogi, Y. Breitbart, H. Korth, and A. Silberschatz. Ensuring transaction atomicity in multidatabase systems. In *ACM Symp. on Principles of Database Systems*, 1992.
- [520] S. Mehrotra, R. Rastogi, H. Korth, and A. Silberschatz. The concurrency control problem in multidatabases: Characteristics and solutions. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1992.
- [521] M. Mehta, R. Agrawal, and J. Rissanen. SLIQ: A fast scalable classifier for data mining. In *Proc. Intl. Conf. on Extending Database Technology*, 1996.
- [522] M. Mehta, V. Soloviev, and D. DeWitt. Batch scheduling in parallel database systems. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1993.
- [523] J. Melton. *Advanced SQL:1999, Understanding Understanding Object-Relational and Other- Advanced Features*. Morgan Kaufmann, 2002.
- [524] J. Melton and A. Sirnon. *Understanding the New SQL: A Complete Guide*. Morgan Kaufmann, 1993.
- [525] J. Melton and A. Sirnon. *SQL:1999, Understanding Relational Language Components*. Morgan Kaufmann, 2002.
- [526] D. Menasce and R. Muntz. Locking and deadlock detection in distributed data bases. *IEEE Transactions on Software Engineering*, 5(3):195–222, 1979.
- [527] A. Mendelzon and T. Milo. Formal models of web queries. In *ACM Symp. on Principles of Database Systems*, 1997.
- [528] A. O. Mendelzon, G. A. Mihaila, and T. Milo. Querying the World Wide Web. *Journal on Digital Libraries*, 1:54–67, 1997.
- [529] R. Meo, G. Psaila, and S. Ceri. A new SQL-like operator for mining association rules. In *Proc. Intl. Conf. on Very Large Databases*, 1996.
- [530] T. Merrett. The extended relational algebra, a basis for query languages. In *Databases*. ed. Shneiderman, Academic Press, 1978.
- [531] T. Merrett. *Relational Information Systems*. Reston Publishing Company, 1983.
- [532] D. Michie, D. Spiegelhalter, and C. Taylor, editors. *Machine Learning, Neural and Statistical Classification*. Ellis Horwood, London, 1994.
- [533] Microsoft. *Microsoft ODBC 3.0 Software Development Kit and Programmer's Reference*. Microsoft Press, 1997.

- (534) K. Mikilineni and S. Suo. An evaluation of relational join algorithms in a pipelined query processing environment. *IEEE Transactions on Software Engineering*, 14(6):838–848, 1988.
- (535) R. Miller, Y. Ioannidis, and R. Ramakrishnan. The use of information capacity in schema integration and translation. In *Proc. Intl. Conf. on Very Large Databases*, 1993.
- T. Milo and D. Suein. Index structures for path expressions. In *ICDT: 7th International Conference on Database Theory*, 1999.
- J. Minker (ed.). *Foundations of Deductive Databases and Logic Programming*. Morgan Kaufmann, 1988.
- [538] T. Minoura and G. Wiederhold. Resilient extended true-copy token scheme for a distributed database. *IEEE Transactions in Software Engineering*, 8(3):173–189, 1982.
- [539] G. Mitchell, U. Dayal, and S. Zdonik. Control of an extensible query optimizer: A planning-based approach. In *Proc. Intl. Conf. on Very Large Databases*, 1993.
- [540] A. Moffat and J. Zobel. Self-indexing inverted files for fast text retrieval. *ACM Transactions on Information Systems*, 14(4):349–379, 1996.
- [541] C. Mohan. ARIES/NT: A recovery method based on write-ahead logging for nested. In *PTOC. Intl. Conf. on Very Large Databases*, 1989.
- [542] C. Mohan. Commit LSN: A novel and simple method for reducing locking and latching in transaction processing systems. In *Proc. Intl. Conf. on Very Large Databases*, 1990.
- [543] C. Mohan. ARIES/LHS: A concurrency control and recovery method using write-ahead logging for linear hashing with separators. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1993.
- [544] C. Mohan, D. Haderle, B. Lindsay, H. Pirahesh, and P. Schwarz. ARIES: a transaction recovery method supporting fine-granularity locking and partial rollbacks using write-ahead logging. *ACM Transactions on Database Systems*, 17(1):94–162, 1992.
- [545] C. Mohan and F. Levine. ARIES/IN: An efficient and high concurrency index manager method using write-ahead logging. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1992.
- [546] C. Mohan and B. Lindsay. Efficient commit protocols for the tree of processes model of distributed transactions. In *ACM SIGACT-SIGOPS Symp. on Principles of Distributed Computing*, 1983.
- [547] C. Mohan, B. Lindsay, and R. Obermarck. Transaction management in the R* distributed database management system. *ACM Transactions on Database Systems*, 11(4):378–396, 1986.
- [548] C. Mohan and I. Narang. Algorithms for creating indexes for very large tables without quiescing updates. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1992.
- [549] K. Morris, J. Naughton, Y. Saraiya, J. Ullman, and A. Van Gelder. YAWN! (Yet Another Window on NAIL!). *Database Engineering*, 6:211–226, 1987.
- [550] A. Motro. Superviews: Virtual integration of multiple databases. *IEEE Transactions on Software Engineering*, 13(7):785–798, 1987.
- [551] A. Motro and P. Buneman. Constructing superviews. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1981.
- [552] R. Mukkamala. Measuring the effect of data distribution and replication models on performance evaluation of distributed database systems. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1989.

- [553] I. Mumick, S. Finkelstein, H. Pirahesh, and R. Ramakrishnan. Magic is relevant. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1990.
- [554] I. Mumick, S. Finkelstein, H. Pirahesh, and R. Ramakrishnan. Magic conditions. *ACM Transactions on Database Systems*, 21(1):107–155, 1996.
- [555] I. Mumick, H. Pirahesh, and R. Ramakrishnan. Duplicates and aggregates in deductive databases. III *Proc. Intl. Conf. on Very Large Databases*, 1990.
- [556] I. Mumick and K. Ross. Noodle: A language for declarative querying in an object-oriented database. In *Intl. Conf. on Deductive and Object-Oriented Databases*, 1993.
- [557] M. Muralikrishna. Improved unnesting algorithms for join aggregate SQL queries. In *Proc. Intl. Conf. on Very Large Databases*, 1992.
- [558] M. Muralikrishna and D. DeWitt. Equi-depth histograms for estimating selectivity factors for multi-dimensional queries. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1988.
- [559] S. Naqvi. Negation as failure for first-order queries. In *ACM Symp. on Principles of Database Systems*, 1986.
- [560] M. Negri, G. Pelagatti, and L. Sbattella. Formal semantics of SQL queries. *ACM Transactions on Database Systems*, 16(3), 1991.
- [561] S. Nestorov, J. Ullman, J. Weiner, and S. Chawathe. Representative objects: Concise representations of semi-structured, hierarchical data. In *Proc. Intl. Conf. on Data Engineering*. IEEE Computer Society, 1997.
- [562] R. T. Ng and J. Han. Efficient and effective clustering methods for spatial data mining. In *Proc. Intl. Conf. on Very Large Databases*, Santiago, Chile, September 1994.
- [563] R. T. Ng, L. V. S. Lakshmanan, J. Han, and A. Pang. Exploratory mining and pruning optimizations of constrained association rules. In *Proc. ACM SIGMOD Intl. Conf. on Management of Data*, pages 13–24. ACM Press, 1998.
- [564] T. Nguyen and V. Srinivasan. Accessing relational databases from the World Wide Web. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1996.
- [565] J. Nievergelt, H. Hinterberger, and K. Sevcik. The Grid File: An adaptable symmetric multikey file structure. *ACM Transactions on Database Systems*, 9(1):38–71, 1984.
- [566] C. Nyberg, T. Barclay, Z. Cvetanovic, J. Gray, and D. Lomet. Alphasort: a cache-sensitive parallel external sort. *VLDB Journal*, 4(4):603–627, 1995.
- [567] R. Obermarck. Global deadlock detection algorithm. *ACM Transactions on Database Systems*, 7(2):187–208, 1981.
- [568] L. O'Callaghan, N. Mishra, A. Meyerson, S. Guha, and R. Iyer. Streaming-data algorithms for high-quality clustering. In *Proc. of the Intl. Conf. on Data Engineering*. IEEE, 2002.
- [569] F. Olken and D. Rotem. Simple random sampling from relational databases. In *Proc. Intl. Conf. on Very Large Databases*, 1986.
- [570] F. Olken and D. Rotem. Maintenance of materialized views of sampling queries. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1992.
- [571] C. Olston, B. T. Loo, and J. Widom. Adaptive precision setting for cached approximate values. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 2001.
- [572] C. Olston and J. Widom. Offering a precision-performance tradeoff for aggregation queries over replicated data. In *Proc. of the Conf. on Very Large Databases*, pages 144–155, 2000.

- [573] C. Olston and T. Widom. Best-effort cache synchronization with source cooperation. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 2002.
- [574] P. O'Neil and E. O'Neil. *Database Principles, Programming, and Performance*. Addison Wesley, 2 edition, 2000.
- [575] P. O'Neil and D. Quass. Improved query performance with variant indexes. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1997.
- [576] B. Ozden, R. Rastogi, and A. Silberschatz. Multimedia support for databases. In *ACM Symp. on Principles of Database Systems*, 1997.
- [577] G. Ozsoyoglu, K. Du, S. Guruswamy, and V.-C. Hou. Processing real-time, non-aggregate queries with time-constraints in case-db. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1992.
- [578] G. Ozsoyoglu, Z. Ozsoyoglu, and V. Matos. Extending relational algebra and relational calculus with set-valued attributes and aggregate functions. *ACM Transactions on Database Systems*, 12(4):566--592, 1987.
- [579] Z. Ozsoyoglu and L.-Y. Yuan. A new normal form for nested relations. *ACM Transactions on Database Systems*, 12(1):111--136, 1987.
- [580] M. Ozsu and P. Valduriez. *Principles of Distributed Database Systems*. PrenticeHall, 1991.
- [581] C. Papadimitriou. The serializability of concurrent database updates. *Journal of the ACM*, 26(4):631--653, 1979.
- [582] C. Papadimitriou. *The Theory of Database Concurrency Control*. Computer Science Press, 1986.
- [583] Y. Papakonstantinou, S. Abiteboul, and H. Garcia-Molina. Object fusion in mediator systems. In *Proc. Intl. Conf. on Very Large Data Bases*, 1996.
- [584] Y. Papakonstantinou, H. Garcia-Molina, and J. Widom. Object exchange across heterogeneous information sources. In *Proc. Intl. Conf. on Data Engineering*, 1995.
- [585] J. Park and A. Segev. Using common subexpressions to optimize multiple queries. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1988.
- [586] J. Patel, J.-B. Yu, K. Tufte, B. Nag, J. Burger, N. Hall, K. Ranasarny, R. Lueder, C. Ellison, J. Kupsch, S. Guo, D. DeWitt, and T. Naughton. Building a scalable geo-spatial DBMS: Technology, implementation, and evaluation. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1997.
- [587] D. Patterson, G. Gibson, and R. Katz. RAID: redundant arrays of inexpensive disks. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1988.
- [588] H.-B. Paul, H.-J. Schek, M. Scholl, G. vVeikurn, and U. Deppisch. Architecture and implementation of the Darmstadt database kernel system. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1987.
- [589] J. Peckham and F. Ilyanski. Semantic data models. *ACM Computing Surveys*, 20(3):153--189, 1988.
- [590] T. Pei and J. Han. Can we push more constraints into frequent pattern mining? In *ACM SIGKDD Conference*, pages 350--354, 2000.
- [591] T. Pei, J. Han, and L. V. S. Lakshmanan. Mining frequent item sets with convertible constraints. In *Proc. Intl. Conf. on Data Engineering (ICDE)*, pages 433--442. IEEE Computer Society, 2001.

- [592] E. Petajan, Y. Jean, D. Lieuwen, and V. Anuparn. DataSpace: An autoillated visu-
alization system for large databases. In *Proc. of SPIE, Visual Data Exploration and
Analysis*, 1997.
- [593] S. Petrov. Finite axiomatization of languages for representation of system properties.
Information Sciences, 47:339-372, 1989.
- [594] G. Piatetsky-Shapiro and C. Cornell. Accurate estimation of the number of tuples
satisfying a condition. In *Proc. ACM SIGMOD Conf. on the Management of Data*,
1984.
- [595] G. Piatetsky-Shapiro and W. J. Frawley, editors. *Knowledge Discovery in Databases*.
AAAI/MIT Press, Menlo Park, CA, 1991.
- [596] H. Pirahesh and J. Hellerstein. Extensible/rule-based query rewrite optimization in
starburst. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1992.
- [597] N. Pitts-Moultis and C. Kirk. *XML black book: Indispensable problem solver*. Corialis
Group, 1998.
- [598] V. Poosala, Y. Ioannidis, P. Haas, and E. Shekita. Improved histograms for selectivity
estimation of range predicates. In *Proc. ACM SIGMOD Conf. on the Management of
Data*, 1996.
- [599] C. Pu. Superdatabases for composition of heterogeneous databases. In *Proc. IEEE Intl.
Conf. on Data Engineering*, 1988.
- [600] C. Pu and A. Leff. Replica control in distributed systems: An asynchronous approach.
In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1991.
- [601] X.-L. Qian and G. Wiederhold. Incremental recomputation of active relational expres-
sions. *IEEE Transactions on Knowledge and Data Engineering*, 3(3):337-341, 1990.
- [602] D. Quass, A. Rajaraman, Y. Sagiv, and J. Ullman. Querying semi-structured heteroge-
neous information. In *Proc. Intl. Conf. on Deductive and Object-Oriented Databases*,
1995.
- [603] J. R. Quinlan. *C4.5: Programs for Machine Learning*. Morgan Kaufman, 1993.
- [604] H. G. M. R. Alonso, D. Barbara. Data caching issues in an information retrieval system.
ACM Transactions on Database Systems, 15(3), 1990.
- [605] The RAIDBook: A source book for RAID technology. The RAID Advisory Board,
<http://www.raid-advisory.com>. North Grafton, MA. Dec. 1998. Sixth Edition.
- [606] D. Rafiei and A. Mendelzon. Similarity-based queries for time series data. In *Proc.
ACM SIGMOD Conf. on the Management of Data*, 1997.
- [607] M. Ramakrishna. An exact probability model for finite hash tables. In *Proc. IEEE Intl.
Conf. on Data Engineering*, 1988.
- [608] M. Ramakrishna and P.-A. Larson. File organization using composite perfect hashing.
ACM Transactions on Database Systems, 14(2):231-263, 1989.
- [609] I. Ramakrishnan, P. Rao, K. Sagonas, T. Swift, and D. Warren. Efficient tabling
mechanisms for logic programs. In *Intl. Conf. on Logic Programming*, 1995.
- [610] R. Ramakrishnan, I. Donjerkovic, A. Ranganathan, K. Beyer, and M. Krishnaprasad.
SRQL: Sorted relational query language. In *PTOC. IEEE Intl. Conf. on Scientific and
Statistical DBMS*, 1998.
- [611] R. Ramakrishnan, D. Srivastava, and S. Sildarshan. Efficient bottom-up evaluation of
logic programs. In *The State of the Art in Computer Systems and Software Engineering*.
ed. J. Vandewalle, Kluwer Academic, 1992.

- [612] R. Ramakrishnan, D. Srivastava, S. Sudarshan, and P. Seshadri. The CORAL: deductive system. *VLDB Journal*, 3(2):161–210, 1994.
- [613] R. Ramakrishnan, S. Stolfo, R. J. Bayardo., and I. Parsa, editors. *Proc. ACM SIGMOD Int. Conference on Knowledge Discovery and Data Mining*. AAAI Press, 2000.
- [614] R. Ramakrishnan and J. Ullman. A survey of deductive database systems. *Journal of Logic Programming*, 23(2):125–149, 1995.
- [615] K. Ramamohanarao. Design overview of the Aditi deductive database system. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1991.
- [616] K. Ramamohanarao, J. Shepherd, and R. Sacks-Davis. Partial-match retrieval for dynamic files using linear hashing with partial expansions. In *Intl. Conf. on Foundations of Data Organization and Algorithms*, 1989.
- [617] V. Raman, B. Raman, and J. M. Hellerstein. Online dynamic reordering for interactive data processing. In *Proc. of the Conf. on Very Large Databases*, pages 709–720. Morgan Kaufmann, 1999.
- [618] S. Rao, A. Badia, and D. Van Gucht. Providing better support for a class of decision support queries. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1996.
- [619] R. Rastogi and K. Shim. Public: A decision tree classifier that integrates building and pruning. In *Proc. Intl. Conf. on Very Large Databases*, 1998.
- [620] D. Reed. Implementing atomic actions on decentralized data. *ACM Transactions on Database Systems*, 1(1):3–23, 1983.
- [621] G. Reese. *Database Programming With JDBC and Java*. O'Reilly & Associates, 1997.
- [622] R. Reiter. A sound and sometimes complete query evaluation algorithm for relational databases with null values. *Journal of the ACM*, 33(2):349–370, 1986.
- [623] E. Rescorla. *SSL and TLS: Designing and Building Secure Systems*. Addison Wesley Professional, 2000.
- [624] A. Reuter. A fast transaction-oriented logging scheme for undo recovery. *IEEE Transactions on Software Engineering*, 6(4):348–356, 1980.
- [625] A. Reuter. Performance analysis of recovery techniques. *ACM Transactions on Database Systems*, 9(4):526–559, 1984.
- [626] E. Riloff and L. Hollaar. Text databases and information retrieval. In *Handbook of Computer Science*. ed. A.B. Tucker, Wiley Press, 1996.
- [627] J. Rissanen. Independent components of relations. *ACM Transactions on Database Systems*, 2(4):317–325, 1977.
- [628] R. Rivest. Partial match retrieval algorithms. *SIAM Journal on Computing*, 5(1):19–50, 1976.
- [629] R. L. Rivest, A. Shamir, and L. M. Adleman. A method for obtaining digital signatures and public-key cryptosystems. *Communications of the ACM*, 21(2):120–126, 1978.
- [630] J. T. Robinson. The KDB tree: A search structure for large multidimensional dynamic indexes. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, 1981.
- [631] J. Hohmer, F. Lescov, and J. Kersey. The Alexander method, a technique for the processing of recursive queries. *New Generation Computing*, 4(3):273–285, 1986.
- [632] D. Rosenkrantz, R. Stearns, and P. Lewis. System level concurrency control for distributed database systems. *ACM Transactions on Database Systems*, 3(2), 1978.
- [633] A. Rosenthal and U. Chakravarthy. Anatomy of a modular multiple query optimizer. In *PTOC. Intl. Conf. on Very Large Databases*, 1988.

- [634] K. Ross and D. Srivastava. Fast coruputation of sparse datacubes. In *Proc. Intl. Conf. on Very Large Databases*, HJ97.
- [635] K. Ross, D. Srivastava, and S. Sudarshan. Materialized view nlaintenance and integrity constraint checking: Trading space for time. In *Proc. ACM SIGMOD Conf. on the Managernent of Data*, 1996.
- [636] J. Rothllie, P. Bernstein, S. Fox, N. Goodnlan, M. HamITler, T. Landers, C. Reeve, D. Shipman, and E. Wong. Introduction to a systeln for distributed databases (SDD -1). *ACM Transactions on Database Systems*, 5(1), 1980.
- [637] J. Rothnie and N. Goodman. An overview of the prelirninary design of SDD -1: A system for distributed data bases. In *Proc. Berkeley Workshop on Dist", "ibuted Data Alanagement and Computer Networks*, 1977.
- [638] N. Roussopoulos, Y. Kotidis, and M. Roussopoulos. Cubetree: Organization of and bulk updates on the data cube. In *PTOC. ACM SIGNIOD Conf. on the Management of Data*, 1997.
- [639] S. Rozen and D. Shasha. Using feature set compromise to automate physical database design. In *Proc. Intl. Conf. on Very LaTge Databases*, 1991.
- [640] J. Rumbaugh, I. Jacobson, and G. Booch. *The Unified Modeling Language Reference Manual (Addison-Wesley Object Technology Series)*. Addison-Wesley, 1998.
- [641] M. Rusinkiewicz, A. Sheth, and G. Karabatis. Specifying interdatabase dependencies in a multidatabase environment. *IEEE ComputeT*, 24(12), 1991.
- [642] D. Sacca and C. Zaniolo. Magic counting methods. In *PTOC. ACM SIGMOD Conf. on the Management of Data*, 1987.
- [643] Y. Sagiv and M. Yannakakis. Equivalence among expressions with the union and difference operators. *Journal of the ACM*, 27(4):633–655, 1980.
- [644] K. Sagonas, T. Swift, and D. Warrell. XSB as an efficient deductive database engine. In *PTOC. ACM SIGMOD Conf. on the Management of Data*, 1994.
- [645] A. Sahuguet, L. Dupont, and T. Nguyen. Kweelt: Querying XivIL in the new millenium. <http://kweelt.sourceforge.net>, Sept 2000.
- [646] G. Salton and M. J. McGill. *IntToduction to Modern Information Retrieval*. McGraw-Hill, 1983.
- [647] B. Salzberg, A. Tsukerman, J. Gray, M. Stewart, S. Uren, and B. Vaughan. Fastsort: A distributed single-input single-output external sort. In *PTOC. ACM SIGMOD Conf. on the Management of Data*, 1990.
- [648] B. J. Salzberg. *Pile StructuTes*. PrenticeHall, 1988.
- [649] H. Salnet. The Quad Tree and related hierarchical data structures. *ACM Computing Surveys*, 16(2), 1984.
- [650] H. Sarnet. *The Design and Analysis of Spatial Data Structures*. Addison-Wesley, 1990.
- [651] J. Sander, M. Ester, H.-P. Kriegel, and X. Xu. Density-based clustering in spatial databases. 1. *of Data Mining and Knowledge Discovery*, 2(2), 1998.
- [652] R. E. Sanders. *ODBC 3.5 Developer's Guide*. McGraw-Hill Series on Data Warehousing and Data Management. McGraw-Hill, 1998.
- [653] S. Sarawagi and M. Stonebraker. Efficient organization of large multidirnsional arrays. In *Proc. IEEE Inil. Conf. on Data Engineering*, 1994.
- [654] S. Sarawagi, S. T'holnas, and R. Agrawal. Integrating mining with relational database systems: Alternatives and inpllications. In *Proc. ACM SIGMOD Intl. Conf. on Man- agernent of Data*, 1998.

- [655] A. Savasere, E. Omiecinski, and S. Navathe. An efficient algorithm for ruining association rules in large databases. In *Proc. Intl. Conf. on Very Large Databases*, 1995.
- [656] P. Schauble. Spider: A multiuser illformlatioll retrieval systetll for semistructured and dyanallic data. In *Proc. ACM 8IGIR Conference on Research and Developrnt in Information Retrieval*, pages 318 - 327, 1993.
- [657] H.-J. Schek, H.-B. Paul, M. Scholl, and G. Weikulll. The DASDBS project: Objects, experiences, and future projects. *IEEE Transactions on Knowledge and Data Engineering*, 2(1), 1990.
- [658] M. Schkolnick. Physical database design techniques. In *NYU Symp. on Database Design*, 1978.
- [659] M. Schkolnick and P. Sorenson. The effects of denormalization on database performance. Technical report, IBM RJ3082, San Jose, CA, 1981.
- [660] G. Schlageter. Optimistic methods for concurrency control in distributed database systems. In *Proc. Intl. Conf. on Very Large Databases*, 1981.
- [661] B. Schneier. *Applied Cryptography: Protocols, Algorithms, and Source Code in C*. John Wiley & Sons, 1995.
- [662] E. Sciore. A complete axiomatization of full join dependencies. *Journal of the ACM*, 29(2):373--393, 1982.
- [663] E. Sciore, M. Siegel, and A. Rosenthal. Using semantic values to facilitate interoperability among heterogeneous information systems. *ACM Transactions on Database Systems*, 19(2):254-290, 1994.
- [664] A. Segev and J. Park. Nlaintaining materialized views in distributed databases. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1989.
- [665] A. Segev and A. Shoshani. Logical rnodeling of ternporal data. *Proc. ACM SIGMOD Conf. on the Management of Data*, 1987.
- [666] P. Selfridge, D. Srivastava, and L. Wilson. IDEA: Interactive data exploration and analysis. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1996.
- [667] P. Selinger and M. Adiba. Access path selections in distributed data base management systems. In *Proc. Intl. Conf. on Databa8es., Brit'ish Computer Society*, 1980.
- [668] P. Selinger, M. Astrahan, D. Charnberlin, R. Lorie, and T. Price. Access path selection in a relational database management system. In *Pl'OC. ACM SIGMOD Conf. on the lvyfanagement of Data*, 1979.
- [669] T. K. Sellis. Multiple query optimization. *AC/VI Transactions on Database Systems*, 13(1):23--52, 1988.
- [670] P. Seshadri, J. Hellerstein, H. Pirahesh, T. Leung, R. Rarnakrishnan, D. Srivastava, P. Stuckey, and S. Sudarshan. Cost-based optirnization for Magic: Algebra and itnplementation. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1996.
- [671] P. Seshadri, M. Livny, and R. Ratnakrishnan. The design and ilnpletuentation of a sequence database system. In *Proc. Intl. Conf. on Very Large Databases*, 1996.
- [672J P. Seshadri, IV!. Livny, and R. Ramakrishnan. The case for enhanced abstract data types. In *Proc. Intl. Conf. on Very Lar:qe Databases*, 1997.
- [67:3] P. Seshadri, H. Pirahesh, and T. Leung. Cornplex query decorrelation. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1996.
- [674] J. Shafer and R. Agrawal. S,PRINT: a scalable parallel classifier for data tlining. In *Proc. Intl. Conf. on Very Large Databases*, 1996.

- [675] J. Shanmugasundaram, U. Fayyad, and P. Bradley. Compressed data cubes for olap aggregate query approximation on continuous dimensions. In *Proc. Intl. Conf. on Knowledge Discovery and Data Mining (KDD)*, 1999.
- [676] J. Shanmugasundaram, J. Kiernan, E. J. Shekita, C. Fan, and J. Funderburk. Querying XML views of relational data. In *Proc. Intl. Conf. on Very Large Data Bases*, 2001.
- [677] L. Shapiro. Join processing in database systems with large main memories. *ACM Transactions on Database Systems*, 11(3):239–264, 1986.
- [678] D. Shasha and N. Goodman. Concurrent search structure algorithms. *ACM Transactions on Database Systems*, 13:53–90, 1988.
- [679] D. Shasha, E. Siruoll, and P. Valduriez. Simple rational guidance for chopping up transactions. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1992.
- [680] H. Shatkay and S. Zdonik. Approximate queries and representations for large data sequences. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1996.
- [681] T. Sheard and D. Sternple. Automatic verification of database transaction safety. *ACM Transactions on Database Systems*, 1989.
- [682] S. Shenoy and Z. Ozsoyoglu. Design and implementation of a semantic query optimizer. *IEEE Transactions on Knowledge and Data Engineering*, 1(3):344–361, 1989.
- [683] P. Shenoy, J. Haritsa, S. Sudarshan, G. Bhalotia, M. Bawa, and D. Shah. Turbocharging vertical mining of large databases. In *Proc. ACM SIGMOD Intl. Conf. on Management of Data*, pages 22–33, May 2000.
- [684] A. Sheth and J. Larson. Federated database systems for managing distributed, heterogeneous, and autonomous databases. *Computing Surveys*, 22(3):183–236, 1990.
- [685] A. Sheth, J. Larson, A. Cornelio, and S. Navathe. A tool for integrating conceptual schemas and user views. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1988.
- [686] A. Shoshani. OLAP and statistical databases: Similarities and differences. In *ACM Symp. on Principles of Database Systems*, 1997.
- [687] A. Shukla, P. Deshpande, J. Naughton, and K. Ramasamy. Storage estimation for multidimensional aggregates in the presence of hierarchies. In *Proc. Intl. Conf. on Very Large Databases*, 1996.
- [688] M. Siegel, E. Sciore, and S. Salveter. A method for automatic rule derivation to support semantic query optimization. *ACM Transactions on Database Systems*, 17(4), 1992.
- [689] A. Silberschatz, H. Korth, and S. Sudarshan. *Database System Concepts (4th ed.)*. McGraw-Hill, 4 edition, 2001.
- [690] E. Simoll, J. Kiernan, and C. de Maindreville. Implementing high-level active rules on top of relational databases. In *Proc. Intl. Conf. on Very Large Databases*, 1992.
- [691] E. Silnoudis, J. Wei, and U. M. Fayyad, editors. *Proc. Intl. Conf. on Knowledge Discovery and Data Mining*. AAAI Press, 1996.
- [692] D. Skeen. Nonblocking commit protocols. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1981.
- [693] J. Srnith and D. Srnith. Database abstractions: Aggregation and generalization. *ACM Transactions on Database Systems*, 1(1):105–133, 1977.
- [694] K. Slnith and M. Winslett. Entity modeling in the MLS relational model. In *Proc. Intl. Conf. on Very Large Databases*, 1992.
- [G95] P. Srnith and M. Barnes. *Files and Databases: An Introduction*. Addison-Wesley, 1987.

- [696] N. Soparkar, H. Korth, and A. Silberschatz. Databases with deadline and contingency constraints. *IEEE Transactions on Knowledge and Data Engineering*, 7(4):552–565, 1995.
- [697] S. Spaccapietra, C. Parent, and Y. Dupont. Model independent assertions for integration of heterogeneous schemas. In *Proc. Intl. Conf. on Very Large Databases*, 1992.
- [698] S. Spaccapietra (ed.). *Entity-Relationship Approach: Ten Years of Experience in Information Modeling*, *Proc. Entity-Relationship Conf.* North-Holland, 1987.
- [699] E. Spertus. ParaSite: Mining structural information on the web. In *Intl. World Wide Web Conference*, 1997.
- [700] R. Srikant and R. Agrawal. Mining generalized association rules. In *Proc. Intl. Conf. on Very Large Databases*, 1995.
- [701] R. Srikant and R. Agrawal. Mining Quantitative Association Rules in Large Relational Tables. In *Proc. ACM SIGMOD Conf. on Management of Data*, 1996.
- [702] R. Srikant and R. Agrawal. Mining Sequential Patterns: Generalizations and Performance Improvements. In *Proc. Intl. Conf. on Extending Database Technology*, 1996.
- [703] R. Srikant, Q. VU, and R. Agrawal. Mining Association Rules with Item Constraints. In *Proc. Intl. Conf. on Knowledge Discovery in Databases and Data Mining*, 1997.
- [704] V. Srinivasan and M. Carey. Performance of B-Tree concurrency control algorithms. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1991.
- [705] D. Srivastava, S. Dar, H. Jagadish, and A. Levy. Answering queries with aggregation using views. In *Proc. Intl. Conf. on Very Large Databases*, 1996.
- [706] D. Srivastava, R. Ramakrishnan, P. Seshadri, and S. Sudarshan. Coral++: Adding object-orientation to a logic database language. In *Proc. Intl. Conf. on Very Large Databases*, 1993.
- [707] J. Srivastava and D. Rotem. Analytical modeling of materialized view maintenance. In *ACM Symp. on Principles of Database Systems*, 1988.
- [708] J. Srivastava, J. Tan, and V. Lum. Tbsam: An access method for efficient processing of statistical queries. *IEEE Transactions on Knowledge and Data Engineering*, 1(4):414–423, 1989.
- [709] D. Stacey. Replication: DB2, Oracle or Sybase? *Database Programming and Design*, pages 42–50, December 1994.
- [710] P. Stachour and B. Thuraisingham. Design of LDV: A multilevel secure relational database management system. *IEEE Transactions on Knowledge and Data Engineering*, 2(2), 1990.
- [711] J. Stankovic and W. Zhao. On real-time transactions. In *Proc. ACM SIGMOD Conf. on the Management of Data Record*, 1988.
- [712] T. Steel. Interim report of the ANSI-SPARC study group. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1975.
- [713] M. Stonebraker. Implementation of integrity constraints and views by query modification. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1975.
- [714] M. Stonebraker. Concurrency control and consistency of multiple copies of data in Distributed Ingres. *IEEE Transactions on Software Engineering*, 5(3), 1979.
- [715] M. Stonebraker. Operating system support for database management. *Communications of the ACM*, 14(7):412–418, 1981.

- [716] M. Stonebraker. Inclusion of new types in relational database systems. In *Proc. IEEE Int'l. Conf. on Data Engineering*, 1986.
- [717] M. Stonebraker. *The INGRES Papers: Anatomy of a Relational Database System*. Addison-Wesley, 1986.
- [718] M. Stonebraker. The design of the Postgres storage system. In *Proc. Intl. Conf. on Very Large Databases*, 1987.
- [719] M. Stonebraker. *Object-relational DBMSs—The Next Great Wave*. Morgan Kaufmann, 1996.
- [720] M. Stonebraker, J. Frew, K. Gardels, and J. Meredith. The Sequoia 2000 storage benchmark. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1993.
- [721] M. Stonebraker and J. Hellerstein (eds). *Readings in Database Systems*. Morgan Kaufmann, 2 edition, 1994.
- [722] M. Stonebraker, A. Jhingran, J. Goh, and S. Potarnianos. On rules, procedures, caching and views in data base systems. In *UCBERL M9036*, 1990.
- [723] M. Stonebraker and G. Kernnitz. The Postgres next-generation database management system. *Communications of the ACM*, 34(10):78--92, 1991.
- [724] B. Subramanian, T. Leung, S. Vandenberg, and S. Zdonik. The AQUA approach to querying lists and trees in object-oriented databases. In *Proc. IEEE Int'l. Conf. on Data Engineering*, 1995.
- [725] W. Sun, Y. Ling, N. Rish, and Y. Deng. An instant and accurate size estimation method for joins and selections in a retrieval-intensive environment. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1993.
- [726] A. Swami and A. Gupta. Optimization of large join queries: Combining heuristics and combinatorial techniques. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1989.
- [727] T. Swift and D. Warren. An abstract machine for SLG resolution: Definite programs. In *Intl. Logic Programming Symposium*, 1994.
- [728] A. Tansel, J. Clifford, S. Chakrabarti, S. Jajodia, A. Segev, and R. Snodgrass. *Temporal Databases: Theory, Design and Implementation*. Benjamin-Cummings, 1993.
- [729] Y. Tay, N. Goodman, and R. Suri. Locking performance in centralized databases. *ACM Transactions on Database Systems*, 10(4):415--462, 1985.
- [730] T. Teorey. *Database Modeling and Design: The E-R Approach*. Morgan Kaufmann, 1990.
- [731] T. Teorey, D.-Q. Yang, and T. Fry. A logical database design methodology for relational databases using the extended entity-relationship model. *ACM Computing Surveys*, 18(2):197--222, 1986.
- [732] R. Thomas. A majority consensus approach to concurrency control for multiple copy databases. *ACM Transactions on Database Systems*, 4(2):180--209, 1979.
- [733] S. A. Thomas. *SSL & TLS Essentials: Securing the Web*. John Wiley & Sons, 2000.
- [734] A. Thomasian. Concurrency control: Methods, performance, and analysis. *ACM Computing Surveys*, 30(1):70--119, 1998.
- [735] A. Thomasian. Two-phase locking performance and its thrashing behavior. *ACM Computing Surveys*, 30(1):70--119, 1998.
- [736] S. Thomas, S. Bodagala, K. Alsabti, and S. Ranka. An efficient algorithm for the incremental updation of association rules in large databases. In *Proc. Intl. Conf. on Knowledge Discovery and Data Mining*. AAAIPress, 1997.

- [737] S. Todd. The Peterlee relational test vehicle. *IBM Systems Journal*, 15(4):285–307, 1976.
- (738) H. Toivonen. Sampling large databases for association rules. In *Proc. Intl. Conf. on Very Large Databases*, 1996.
- [739] TP Performance Council. TPC Benchmark D: Standard specification, rev. 1.2. Technical report, <http://www.tpc.org/dspec.html>, 1996.
- [740] I. Traiger, J. Gray, C. Galtieri, and B. Lindsay. Transactions and consistency in distributed database systems. *ACM Transactions on Database Systems*, 25(9), 1982.
- [741] M. Tsangaris and J. Naughton. On the performance of object clustering techniques. In *Proc. ACVI SIGMOD Conf. on the Management of Data*, 1992.
- [742] D.-M. Tsou and P. Fischer. Decomposition of a relation scheme into Boyce-Codd normal form. *SIGACT News*, 14(3):23–29, 1982.
- [743] D. Tsur, J. D. Ullman, S. Abiteboul, C. Clifton, R. Motwani, S. Nestorov, and A. Rosenthal. Query flocks: A generalization of association-rule mining. In *Proc. ACM SIGMOD Conf. on Management of Data*, pages 1–12, 1998.
- [744] A. Tucker (ed.). *Computer Science and Engineering Handbook*. CRC Press, 1996.
- [745] J. W. Thkey. *Exploratory Data Analysis*. Addison-Wesley, 1977.
- [746] J. Ullman. The U.R. strikes back. In *ACM Symp. on Principles of Database Systems*, 1982.
- [747] J. Ullman. *Principles of Database and Knowledgebase Systems*, Vols. 1 and 2. Computer Science Press, 1989.
- [748] J. Ullman. Information integration using logical views. In *Intl. Conf. on Database Theory*, 1997.
- (749) S. Urban and L. Delcambre. An analysis of the structural, dynamic, and temporal aspects of semantic data models. In *Proc. IEEE Intl. Conf. on Data Engineering*, 1986.
- [750] G. Valentin, M. Zuliani, D. C. Zilio, G. M. Lohman, and A. Skelley. Db2 advisor: An optimizer smart enough to recommend its own indexes. In *Proc. Intl. Conf. on Data Engineering (ICDE)*, pages 101–110. IEEE Computer Society, 2000.
- [751] M. Van Emden and R. Kowalski. The semantics of predicate logic as a programming language. *Journal of the ACM*, 23(4):733–742, 1976.
- [752] A. Van Gelder. Negation as failure using tight derivations for general logic programs. In J. Minker, editor, *Foundations of Deductive Databases and Logic Programming*. Morgan Kaufmann, 1988.
- [753] C. J. van Rijsbergen. *Information Retrieval*. Butterworths, London, United Kingdom, 1990.
- [754] M. Vardi. Incomplete information and default reasoning. In *ACM Symp. on Principles of Database Systems*, 1986.
- [755] M. Vardi. Fundamentals of dependency theory. In *Trends in Theoretical Computer Science*. ed. E. Borger, Computer Science Press, 1987.
- [756] L. Vieille. Recursive axioms in deductive databases: The query-subquery approach. In *Intl. Conf. on Expert Database Systems*, 1986.
- (757) L. Vieille. From QSQ towards QoSAQ: global optimization of recursive queries. In *Intl. Conf. on Expert Database Systems*, 1988.
- [758] L. Vieille, P. Bayer, V. Kuchenhoff, and A. Lefebvre. EKS-VI, a short overview. In *AAAI-90 Workshop on Knowledge Base Management Systems*, 1990.

- [759] J. S. Vitter and M. Wang. Approximate computation of multidimensional aggregates of sparse data using wavelets. In *Proc. ACM SIGMOD Conf. on the Management of Data*, pages 193-204. ACM Press, 1999.
- [760] G. von Bultzingsloewen. Translating and optimizing SQL queries having aggregates. In *Proc. Intl. Conf. on Very Large Databases*, 1987.
- [761] G. von Bultzingsloewen, K. Dittrich, C. Iochpe, R.-P. Liedtke, P. Lockemaun, and M. Schryro. Kardamom-----A dataflow database machine for real-time applications. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1988.
- [762] G. Vossen. *Data Models, Database Languages and Database Management Systems*. Addison-Wesley, 1991.
- [763] N. Wade. Citation analysis: A new tool for science administrators. *Science*, 188(4183):429-432, 1975.
- [764] R. Wagner. Indexing design considerations. *IBM Systems Journal*, 12(4):351-367, 1973.
- [765] X. Wang, S. Jajodia, and V. Subrahmanian. Temporal modules: An approach toward federated temporal databases. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1993.
- [766] K. Wang and H. Liu. Schema discovery for semistructured data. In *Third International Conference on Knowledge Discovery and Data Mining (KDD -97)*, pages 271-274, 1997.
- [767] R. Weber, H. Sehek, and S. Blott. A quantitative analysis and performance study for similarity-search methods in high-dimensional spaces. In *Proc. Intl. Conf. on Very Large Data Bases*, 1998.
- [768] G. Weddell. Reasoning about functional dependencies generalized for semantic data models. *ACM Transactions on Database Systems*, 17(1), 1992.
- [769] W. Weihl. The impact of recovery on concurrency control. In *ACM Symp. on Principles of Database Systems*, 1989.
- [770] G. Weikum and G. Vossen. *Transactional Information Systems*. Morgan Kaufmann, 2001.
- [771] R. Weiss, B. V. Iez, M. A. Sheldon, C. Manprelpre, P. Szilagyi, A. Duda, and D. K. Gifford. HyPursuit: A hierarchical network search engine that exploits content-link hypertext clustering. In *Proc. ACM Conf. on Hypertext*, 1996.
- [772] C. White. Let the replication battle begin. In *Database Programming and Design*, pages 21-24, May 1994.
- [773] S. White, M. Fisher, R. Cattell, G. Hanlilton, and M. Hapner. *JDBC API Tutorial and Reference: Universal Data Access for the Java 2 Platform*. Addison-Wesley, 2 edition, 1999.
- [774] J. Widorn and S. Ceri. *Active Database Systems*. Morgan Kaufmann, 1996.
- [775] G. Wiederhold. *Database Design (2nd ed.)*. McGraw-Hill, 1983.
- [776] G. Wiederhold, S. Kaplan, and D. Sagalowicz. Physical database design research at Stanford. *IEEE Database Engineering*, 1:117-119, 1983.
- [777] R. Williams, D. Daniels, L. Haas, G. Lapis, B. Lindsay, P. Ng, R. Obernarck, P. Selinger, A. Walker, P. Wilms, and R. Yost. R*: An overview of the architecture. Technical report, IBM RJ3325, San Jose, CA, 1981.
- [778] M. S. Winslett. A model-based approach to updating databases with incomplete information. *ACM Transactions on Database Systems*, 13(2):167-196, 1988.

- [779] G. vViorkowski and D. Kull. *DB2: Design and Development Guide (3rd ed.)*. Addison-Wesley, 1992.
- [780] I. H. Witten, A. Moffat, and T. C. Bell. *Managing Gigabytes: Compressing and Indexing Documents and Images*. Van Nostrand Reinhold, 1994.
- [781] I. H. Witten and E. Frank. *Data Mining: Practical Machine Learning Tools and Techniques with Java Implementations*. Morgan Kaufmann Publishers, 1999.
- [782] O. Wolfson, A. Sistla, B. Xu, J. Zhou, and S. Chamberlain. Domino: Databases for moving objects tracking. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, 1999.
- [783] Y. Yang and R. Miller. Association rules over interval data. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1997.
- [784] K. Youssefi and E. Wong. Query processing in a relational database management system. In *Proc. Intl. Conf. on Very Large Databases*, 1979.
- [785] C. Yu and C. Chang. Distributed query processing. *ACM Computing Surveys*, 16(4):399–433, 1984.
- [786] O. R. Zaiane, M. El-Hajj, and P. Lu. Fast Parallel Association Rule Mining Without Candidacy Generation. In *Proc. IEEE Intl. Conf. on Data Mining (ICDM)*, 2001.
- [787] M. J. Zaki. Scalable algorithms for association mining. In *IEEE Transactions on Knowledge and Data Engineering*, volume 12, pages 372–390, May/June 2000.
- [788] M. J. Zaki and C.-T. Ho, editors. *Large-Scale Parallel Data Mining*. Springer Verlag, 2000.
- [789] C. Zaniolo. Analysis and design of relational schemata. Technical report, Ph.D. Thesis, UCLA, TR UCLA-ENG-7669, 1976.
- [790] C. Zaniolo. Database relations with null values. *Journal of Computer and System Sciences*, 28(1):142–166, 1984.
- [791] C. Zaniolo. The database language GEM. In *Readings in Object-Oriented Databases*. eds. S.B. Zdonik and D. Maier, Morgan Kaufmann, 1990.
- [792] C. Zaniolo. Active database rules with transaction-conscious stable-model semantics. In *Intl. Conf. on Deductive and Object-Oriented Databases*, 1996.
- [793] C. Zaniolo, N. Arni, and K. Ong. Negation and aggregates in recursive rules: the LDL++ approach. In *Intl. Conf. on Deductive and Object-Oriented Databases*, 1993.
- [794] C. Zaniolo, S. Ceri, C. Faloutsos, R. Snodgrass, V. Subrahmanian, and R. Zicari. *Advanced Database Systems*. Morgan Kaufmann, 1997.
- [795] S. Zdonik, U. Cetintemel, M. Cherniack, C. Convey, S. Lee, G. Seidman, M. Stonebraker, N. Tatbul, and D. Carney. Monitoring streams—A new class of data management applications. In *Proc. Intl. Conf. on Very Large Data Bases*, 2002.
- [796] S. Zdonik and D. Maier (eds.). *Readings in Object-Oriented Databases*. Morgan Kaufmann, 1990.
- [797] A. Zhang, M. Nodine, B. Bhargava, and O. Bukhres. Ensuring relaxed atomicity for flexible transactions in multidatabase systems. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1994.
- [798] T. Zhang, R. Hamakrishnan, and M. Livny. BIRCH: an efficient data clustering method for very large databases. In *Proc. ACM SIGMOD Conf. on Management of Data*, 1996.
- [799] Y. Zhao, P. Deshpande, J. F. Naughton, and A. Shukla. Simultaneous optimization and evaluation of multiple dimensional queries. In *Proc. ACM SIGMOD Int. Conf. on Management of Data*, 1998.

- [800] Y. Zhuge, H. Garcia-Molina, J. Hanuner, and .l. Widom. View maintenance in a warehousing enviroIllnent. In *Proc. ACM SIGMOD Conf. on the Management of Data*, 1995.
- [801J M. M. Zloof. Query-by-exaIllple: a database language. *IBM Systems Journal*, 16(4):324--343, 1977.
- [802J J. Zobel, ,A. Moffat, and K. RalnarlIohanarao. Inverted files versus signature files for text indexing. *ACM Transactions on Database System,s*, 23, 1998.
- [803] J. Zobel, A. Moffat, and R. Sacks-Davis. An efficient indexing technique for full text databases. In *Froc. Intl. Co'nj. on Very Large Databases., Morgan Kaufman pubs. (San Francisco, CAy is] Vancou'veT*, 1992.
- [804] U. Zukowski and B. Freitag. The deductive database system LOLA. In *Proc. Intl. Conf. on Log'ic Programming and N'on-AI0notonic Reasoning*, 1997.

AUTHOR INDEX

- Abbott, R., 578, 1005, 1001
 Abdali, K., 270, 1015
 Abdellatif, A., 771, 1029
 Abiteboul, S., 24, 98, 648, 816, 844, 925, 967, 1005, 1030, 1033, 1041, 1001
 Abounaga, A., 967, 1005
 Acharya, S., 888, 1005
 Achyutuni, K.J., 578, 1005
 Ackaouy, E., xxix
 Adali, S., 771, 1005
 Adiba, M.E., 771, 1005, 1038
 Adleman, L.W., 722, 1036
 Adya, A., 815, 1029
 Agarwal, R.C., 924, 1005
 Agarwal, S., 887, 1005
 Aggarwal, C.C., 924, 1005
 Agrawal, D., 578, 771, 1005
 Agrawal, R., 181, 602, 815, 887, 924, 925, 1005-1006, 1008, 1024, 1030, 1037-1039
 Ahad, R., 516, 1006
 Ahlberg, C., 1006, 1001
 Ahmed, R., 1026, 1001
 Aho, A.V., 303, 516, 648, 1006
 Aiken, A., 181, 1006, 1001
 Ailamaki, A., 1006
 Alameldeen, A.R., 967, 1005
 Albert, J.A., xxxi, 1026, 1001
 Alon, N., 887
 Alonso, R., 966, 1035
 Alsabti, K., 925, 1041
 Anupam, V., 1034, 1001
 Anwar, E., 181, 1007
 Apt, K.R., 845, 1007
 Armstrong, W.W., 648, 1007
 Arni, N., 816, 1044
 Arocena, G., 967, 1007
 Asgarian, M., 691, 1012
 Astrahan, M.M., 98, 180, 516, 1007, 1012, 1013, 1038
 Atkinson, M.P., 815, 816, 1007
 Attar, R., 771, 1007
 Atzeni, P., 24, 98, 648, 8Ui, 967, 1007
 Avnur, R., 888, 1007
 Babcock, B., 1007
 Babu, S., 888, 1007
 Badal, D.Z., 98, 1007
 Badia, A., 129, 887, 1007, 1035
 Badrinath, B.R., 1025, 1001
 Baeza-Yates, R., 966, 1019
 Bailey, P., 1007
 Balbin, I., 845, 1008
 Ballou, N., 815, 1026
 Balsters, H., xxxi
 Bancelhon, F., 99, 816, 845, 1008
 BapaRao, K.V., 516, 1006
 Baralis, E., 181, 1008
 Barbara, D., 771, 888, 966, 1008, 1020, 1035
 Barclay, T., 438, 1033
 Barnes, N.L.G., 303, 1039
 Barnett, J.R., 337, 1008
 Barquin, R., 887, 1008
 Batini, C., 55-56, 1008
 Batory, D.S., 516, 1008, 1026
 Baugsto, B.A.W., 438, 1008
 Baum, M.S., 722, 1019
 Bawa, M., 924, 1038
 Bayardo, R.J., 924-925, 1008, 1035
 Bayer, P., 844, 1042
 Bayer, R., 369, 1008
 Beck, M., 438, 1008
 Beckmann, N., 991, 1008
 Beech, D., 815, 1018
 Beeri, C., 648, 816, 845, 1006, 1009
 Bektas, H., xxix
 Bell, D., 771, 1009
 Bell, T.E., 966, 1043
 Bentley, J.L., 369, 1009
 Berchtold, S., 991, 1009
 Bergstein, P., xxxii
 Bernstein, A.J., 24, 771, 1027-1028
 Bernstein, P.A., 99, 548, 576, 578, 648, 771, 1007, 1009, 1010, 1015, 1036
 Beyer, K.S., 887, 991, 1010, 1029, 1035, 1038
 Bhalotia, G., 924, 1038
 Bhargava, K., xxxii, 771, HHO
 BiJiris, A., 337, 1010
 Biskup, J., 56, 648, UHO
 Bitton, J., 438, 477, 1008, 1010
 Blajr, H., 845, 1007
 Blakeley, J.A., 887, 1010
 Blanchard, L., xxx
 Blasgen, M.W., 98, 477, 602, 1007, 1010, 1012, 1022
 Blaustein, B.T., 99, 1009
 Blott, S., 991, 1043
 Bodagala, S., 925, 1041
 Bohannon, P., 967, 1010, 1026, 1001
 Bohm, C., 991, 1009
 Bonaparte, N., 847
 Bonnet, P., 1010
 Booe, G., 56, 1010, 1036
 Borat, H., 477, 516, 1027
 Boroclin, A., 966, 1010
 Bosworth, A., 887, 1022
 Boyce, R.F., 180, 1010
 Bradley, P.S., 887, 1010, 1038, 925
 Bratbergsengen, K., 477, 1010
 Breiman, L., 925, 1010
 Breitbart, Y., 771, 1010-1011, 1030
 Brin, S., 924, 966, 1011
 Brinkhoff, T., 991, 1011
 Brown, K.P., 337, 1011
 Bruno, N., 888, 1011
 Bry, F., 99, 845, 1011
 Bukhres, O.A., 771, 1017
 Buneman, A.P., 56, 181, 815-816, 967, 1007, 1011, 1032
 Bunker, R., 477, 1021
 Burdick, D., 924, 1011
 Burger, J., 1034, IDOI
 Burke, E., 422
 Cabibbo, L., 816, 1007
 Cai, L., xxxi
 Calimlim, M., 924, 1011
 Campbell, D., xxxi
 Candan, K.S., 771, 1005
 Carey, M.J., xxix, xxxi-xxxii, 337, 578, 691, 771, 815, 816, 888, 967, 1004, 1006, 1011-1012, 1019, 1022-1023, 1039
 Carney, D., 888, 1044
 Carroll, L., 440
 Casanova, M.A., 56, 99, 1012, 1019

- Castano, S., 722, 1012
 Castro, M., 815, 1029
 Cate, H.P., 815, 1018
 Cattell, R.G.G., 219, 816, 1012, 1023, 1043
 Ceri, S., 55, 99, 181, 771, 816, 844, 925, 1008, 1012, 1031, 1043-1044, 1001
 Cesarini, F., 691, 1012
 Cetintemel, U., 888, 1044
 Chakravarthy, U.S., 181, 516, 578, 1007, 1012, 1024, 1036
 Chamberlain, S., 991, 1043
 Chamberlin, D.D., 98-99, 180-181, 516, 816, 967, 1007, 1010-1013, 1017, 1038
 Chan, M.C., 771
 Chan, P., 924, 1025
 Chandra, A.K., 516, 845, 1013
 Chandy, M.K., 771, 1013
 Chang, C.C., 771, 1013, 1043
 Chang, D., 270, 1013
 Chang, S.K., 771
 Chang, W., 815, 1022
 Chanliau, M., xxxi
 Chao, D., xxxi
 Charikar, M., 888, 1013
 Chatziantoniou, D., 887, 1013
 Chaudhuri, S., 691, 816, 887-888, 924, 1011, 1013
 Chawathe, S., 967, 1032
 Cheiney, J.P., 477, 1013
 Chen, C.M., 337, 516, 1013
 Chen, G., 1029, 1001
 Chen, H., xxxi
 Chen, J., 1006, 1001
 Chen, P.M., 337, 1014
 Chen, P.P.S., 1014
 Chen, Y., 887, 1014
 Cheng, W.H., 771
 Cherniack, M., 888, 1044
 Cheung, D.J., 925, 1014
 Childs, D.L., 98, 1014
 Chimenti, D., 844, 1014
 Chin, F.Y., 722, 1014
 Chisholm, K., 1007
 Chiu, D.W., 771, 1009
 Chiueh, T-C., 966, 1014
 Chomicki, J., 99, 1014
 Chou, H., 337, 815, 1014, 1016
 Chow, E.C., 815, 1018
 Christodoulakis, S., 516, 966, 1025
 Chrysanthis, P.K., 548, 1014
 Chu, F., 516, 1014
 Chu, P., 516, 1030
 Churchill, W., 992
 Civelek, F.N., 56, 1014
 Clarke, E.M., 99, 1009
 Clemons, E.K., 181, 1011
 Clifford, J., 1041, 1001
 Clifton, C., 925, 1041
 Cochrane, R.J., 181, 1014
 Cockshott, P., 1007
 Codd, E.F., 98, 129, 648, 887, 1014-1015
 Colby, L.S., 887, 1015
 Collier, R., 25
 Comer, D., 369, 1015
 Connell, C., 516
 Connolly, D., 270, 1015
 Connors, T., 815, 1018
 Convent, B., 56, 1010
 Convey, C., 888, 1044
 Cooper, B., 967, 1015
 Cooper, S., 477, 1021
 Copeland, D., 815, 1015
 Cornelio, A., 1039, 56
 Cornell, C., 516, 1034
 Cornell, G., 270, 1015
 Cortes, C., 888
 Cosmadakis, S.S., 99
 Cristian, F., 771, 1017
 Cristodoulakis, S., 1018
 Cvetanovic, Z., 438, 1033
 Dadam, P., 337, 816, 1028
 Daemen, J., 722, 1015
 Daniels, D., 771, 1043
 Dar, S., 887, 1040
 Das, G., 888, 1013
 Datal', M., 888, 1007
 Date, C.J., 24, 98-99, 637, 648, 1015
 Davidson, S., 967, 1011
 Davis, J.W., 815, 1018
 Davis, K.C., xxxi
 Dayal, D., 99, 181, 516, 648, 771, 887, 1010, 1013, 1015, 10301031
 Day, M., 815, 1029
 De Antonellis, V., 24, 98, 648, 1007
 De Maindreville, C., 181, 1039
 DeBono, E., 304
 DeBra, P., 648, 1015
 Deep, J., 270, 1015
 Delcambre, L.M.L., xxxi, 56, 1042
 Delobel, C., 648, 816, 1008, 1015
 Deng, Y., 516, 1041
 Denning, D.E., 722, 1015, 1029
 Deppisch, U., 3:37, 1034
 Derr, M., 1016
 Derrett, N., 815, 1018
 Derstad, J., 267
 Deshpande, A., 816, 1016
 Deshpande, P., 887, 1005, 1016, 1039, 1044
 Deutsch, A., 967, 1016
 Deux, O., 815, 1016
 DeWitt, D.J., xxviii, 337, 438, 477, 516, 602, 691, 770-771, 815-816, 1004, 1006, 1010-1012, 1014, 1016, 1021, 1024-1025, 1030, 1032, 1034, 1001
 Diaz, O., 181, 1016
 Dickens, C., 605
 Dietrich, S.W., 845, 887, 1016, 1023
 Diffie, W., 722, 1016
 Dimino, L., xxxi
 Dittrich, K.R., 816, 1042, 1001
 Dogac, A., 56, 771, 1014, 1023
 Domingos, P., 925, 1016
 Domingos, R., 925, 1024
 Dong, G., 887, 1014
 Donjerkovic, D., xxix, 516, 887-888, 1016, 1029, 1035, 1001
 Donne, J., 726
 Doole, D., 816, 1011
 Doraiswamy, S., xxxi
 Doyle, A.C., 773
 Dubes, R., 925
 Dubes, R.C., 1016, 1025
 Du, K., 1033, 1001
 Du, W., 771, 1016
 Duda, A., 966, 1043
 DuMouchel, W., 888, 1008
 Dupont, L., 967, 1037
 Dupont, Y., 56, 1039
 Duppel, N., 477, 1016
 Eaglin, R., xxxii
 Edelstein, H., 771, 887, 1008, 1017
 Effelsberg, W., 337, 1017
 Eich, M.H., xxxi, 602, 1017
 Eisenberg, A., 180, 816, 1017
 El Abbadi, A., 578, 771, 1005, 1017
 El-Hajj, M., 924, 1043
 Ellis, C.8., 578, 1017
 Ellman, C., 1034, 1001
 Elmagarmid, A.K., 771, 1000, 1015-1017
 Elmasri, R., 24, 56, 1017
 Epstein, R., 477, 771, 1017
 Erbe, R., 337, 816, 1028
 Ester, M., 925, 1017, 1037
 Eswaran, K.P., 98, 180181, 477, 548, 1007, uno, 1013, 1017

- Fagin, R., xxix, 390, 637, 648, 1009, 1015, 1017--1018
- Faloutsos, C., 181, 337, 369, 816, 844, 888, 925, 966, 991, 1008, 1018, 1027, 1044, 1001
- Fan, C., 967, 1038
- Fang, M., 924, 1018
- Fandemay, P., 477, 1013
- Fayyad, U.M., 887, 924--925, 1006, 1010, 1018, 1038--1039
- Fendrich, J., xxxii
- Fernandez, M., 967, 1016, 1018
- Finkelstein, S.J., 516, 691, 845, 1018, 1032
- Fischer, C.N., xxx
- Fischer, P.C., 648, 1025, 1041
- Fisher, K., 888
- Fisher, M., 219, 1023, 1043
- Fishman, D.H., 815, 1018
- Fitzgerald, E., 817
- Fleming, C.C., 691, 1019
- Flisakowski, S., xxix--xxx
- Florescu, D., 967, 1012--1013, 1016, 1018-1019
- Ford, W., 722, 1019
- Fotouhi, F., 477, 1019
- :Fowler, M., 56, 1019
- Fox, S., 771, 1036
- Frakes, W.B., 966, 1019
- Franaszek, P.A., 578, 1019
- Franaszek, P.A., 1019
- Frank, E., 924, 1043
- Franklin, M.J., 771, 815816, 967, 1011, 1015, 1019
- :Fraternali, P., 181, 1012, 1019
- Frawley, W.J., 924, 1034
- Freeston, M.W., 991, 1019
- Freire, J., 967, 1010
- Freitag, B., 844, 1044
- French, J., 1020
- Frew, J., 691, 1040
- Freytag, J.C., 516, 1019
- Friedman, J.B., 369, 924 925, 1009---1010, 102:3
- Friesen, O., 816, 1019
- Fry, J.P., 24, 56, 99, 1019, 1041
- Fuchs, IV., 270, 1028
- Fu, Y., 925, 1023
- Fugini, M.G., 722, 1012
- Fuhr, N., 966, 1019
- Fukuda, T., 924, 1019
- Funderburk, J., 967, 1038
- Furtado, A.I., 99, 1012, 1019
- Fushimi, S., 477, 1019
- GacHa, S., 1041, 10CH
- Gaede, V., 991, 1020
- Gallaire, H., 98--99, 648, 844, 1020
- Galtieri, C.A., 602, 771, 1028, 1041
- Gamboa, R., 844, 1014
- Ganguly, S., 771, 1020
- Ganski, R.A., 516, 1020
- Ganti, V., 925, 1020
- Garcia-Molina, H., 24, 578, 771, 887, 924, 966---967, 1005, 1010, 1018, 1020--1021, 1033-1035, 1044, 1001
- Gardels, K., 691, 1040
- Garfield, E., 966, 1020
- Garg, A.K., 390, 1020
- Garza, J.F., 337, 815, 1008, 1026
- Gehani, N.H., 181, 815, 1006
- Gehrke, J.E., 691, 888, 924-925, 1006, 1011--1012, 1020
- Gerber, R.H., 477, 770, 1016
- Ghemawat, S., 815, 1029
- Ghosh, S.P., 303, 1020
- Gibbons, P.B., 887--888, 1005, 1021
- Gibson, D., 925, 966, 1021
- Gibson, G.A., 337, 1014, 1021, 1034
- Gifford, D.K., 771, 1021, 1043
- Gifford, K., 966
- Gilbert, A.C., 888
- Gionis, A., 888
- Goh, J., 181, 1040
- Goldfarb, C.F., 270, 1021
- Goldman, R., 967, 1021, 1030
- Goldstein, J., xxxi, 991, 1010, 1021
- Goldweber, M., xxix
- Goodman, N., 576, 578, 771, 1007, 1009, 1036, 1038, 1041
- Gopalan, H., xxxi
- Gotlieb, C.C., 390, 1020
- Gottlob, G., 844, 1012
- Graefe, G., xxxi, 4:38, 477, 51G, 770--771, 815, 1011, 1016, 1021, 1028
- Graham, M.H., 648, 1021
- Grahne, G., 98, 1021
- Grant, J., 516, 1012
- Gravano, L., 888, 966, 1011, 1021
- Gray, J.N., 98, 438, 548, 602, 691, 770--771, 887, 1000, 1007, 1012, 1016--1017, 1021-1022, 1028, 1033, 1037, 1041
- Gray, P.M.D., 24, 181, 1016, 1022
- Greenwald, M., 887, 1022
- Greipsland, J.F., 438, 1008
- Griffin, T., 887, 1015
- Griffiths, P.P., 98, 180, 602, 722, 1007, 1013, 1022
- Grimson, J., 771, 1009
- Grinstein, G., 1022, 1001
- Grosky, W., xxxi
- Gruber, R., 815, 1029
- Guenther, O., 991, 1020
- Guha, S., 888, 925, 1022, 1033
- Gunopulos, D., 924--925, 1006, 1008, 1022
- Guo, S., 1034, 1001
- Gupta, A., 516, 887, 1005, 1022, 1041
- Guruswamy, S., 1033, 1001
- Guttman, A., 991, 1022
- Gyssens, M., 129, 1007
- Haas, L.M., 771, 815, 1013, 1022, 1043
- Haas, P.J., 516, 888, 1022---1023, 1034
- Haber, E., xxx
- Haderle, D., 602, 771, 1031
- Hadzilacos, V., 576, 578, 1009
- Haerder, T., 337, 602, 1017, 1023
- Haight, D.M., 815, 1011
- Haines, M., xxix
- Halici, U., 771, 1023
- Hall, M., 270, 1023
- Hall, N.E., 815, 1011, 1034, 1001
- Hall, P.A.V., 477, 1023
- Halpern, J.Y., 516, 1014
- Hamilton, G., 219, 1023, 1043
- Hammer, J., xxxi, 887, 1044
- Hammer, M., 98, 771, 1023, 1036
- Han, J., 887, 924--925, 1014, 1023, 1028, 1032, 1034
- Hand, D.J., 924925, 1023, 1030
- Hanson, E.N., 181, 887, 1023
- Hapner, M., 219, 1043
- Barel, D., 845, 1013
- Harinarayan, V., 887, 1023
- Haritsa, J., 578,924,1023, 1038
- Harkey, D., 270, 101:3
- Harrington, J., xxx
- Harris, S., xxix
- Harrison, J., 887, 1023
- Hasan, W., 771, 1020
- Hass, P.or., 888, 1008
- Hastie, T., 924, 1023
- Hearst, M., xxxii

- Heckerman, D., 924, 1014, 1023, 1041
 Heckman, M., 722, 1029
 Helland, P., 771
 Hellerstein, J.M., xxix, 181, 516, 772, 816, 845, 888, 967, 991, 1006-1008, 1022, 1024, 1027, 1034, 1035, 1038, 1040, 773
 Hellman, M.E., 722, 1016
 Heilschen, L.J., 99, 1030
 Heytens, M.L., 477, 770, 1016
 Hidber, C., 925, 1024
 Hill, M.D., 1006
 Hillebrand, G., 967, 1011
 Himmeroeder, R., 967, 1024
 Hinterberger, H., 991, 1033
 Hjaltason, G.R., 967, 1015
 Hoch, C.G., 815, 1018
 Ho, C-T., 887, 924, 1024, 1044
 Holfelder, P., 270, 1015
 Hollaar, L.A., 966, 1036, 1001
 Holzner, S., 270, 1024
 Honeyman, P., 648, 1009
 Hong, D., 578, 1024
 Hong, W., 771, 1024
 Hopcroft, J.E., 303, 1006
 Hou, W-C., 516, 1024, 1033, 1001
 Howard, J.H., 648, 1009
 Hsiao, H., 771, 1024
 Hsu, C., xxxii
 Huang, .1., 578, 1024
 Huang, L., 966, 1014
 Huang, W., xxix
 Huang, Y., 771, 1024, 1001
 Hull, R., 24, 56, 98, 648, 816, 844, 1005, 1024, 1001
 Hulten, G., 925, 1016, 1024
 Hunter, J., 270, 1024
 Imielinski, T., 98, 924, 1006, 1024-1025, 1001
 Joannidis, Y.E., xxix, 56, 516, 888, 1008, 1025, 1031, 1034
 Iochpe, C., 1042, 10CH
 Ives, Z., 967, 1012
 Jacobson, I., 56, 1010, 1036
 Jacobsson, H., xxxi
 Jagadish, H.V., 337, 369, 887-888, 991, 1008, 1018, 1025, 1027, 1040, 1001
 Jain, A.K., 925, 1016, 1025
 Jadodia, S., 722, 771, 1025, 1041-1042, 1001
 .larke, M., 516, 1025
 Jean, Y., 1034, 1001
 Jeffers, R., 578, 1005
 .Jhingran, A., 181, 1040
 Jing, J., 771, 1017
 Johnson, T., 578, 888, 1008, 1024
 Jones, K.S., 966, 1025
 Jonsson, B.T., 771, 1019
 Jou, J.H., 648, 1025
 Kabra, N., 516, 1025, 1034, 1001
 Kambaya,shi, Y., 771, 1025
 Kamber, M., 924, 1023
 Kane, S., xxxii
 Kanellakis, P.C., 98, 648, 816, 1005, 1008, 1025, 1001
 Kang, J., 967, 1018
 Kang, Y.C., 516, 1025
 Kaplan, S.J., 1043
 Karabatis, G., 771, 1036, 1001
 Kargupta, H., 924, 1025
 Katz, R.H., 337, 477, 1014, 1016, 1034
 Kaufman, L., 925, 1025
 Kaushik, R., xxxii, 967, 1026, 926
 Kawaguchi, A., xxxii, 887, 1015
 Keats, J., 130
 Kedem, Z.Ivl., 924, 1028
 Keim, D.A., 1026, 1001
 Keller, A.M., 99, 1026
 Kemnitz, G., 815, 1040
 Kemper, A.A., 337, 816, 1028
 Kent, W., 24, 616, 815, 1018, 1026, 1001
 Kerisit, J.M., 845, 1036
 Kerschberg, L., 24, 1026
 Ketabchi, M.A., 1026, 1001
 Khanna, S., 887, 1022
 Khardon, R., 924, 1022
 Khayyam, O., 817
 Khoshafian, S., 816, 1008
 Kiernan, J., 181, 967, 1038-1039
 Kiessling, W., 516, 1026
 Kifer, M., 24, xxix, 816, 845, 1026, 1028
 Kimball, H., 887, 1026
 Kim, W., 516, 771, 815, 816, 1017, 1026
 Kimmel, W., xxx
 King, J.J., 516, 1026
 King, R., 56, 1024
 King, W.F., 98, 1007, 1012
 Kirk, C., 270, 1034
 Kitsuregawa, M., 477, 1019
 Kleinberg, J.M., 925, 966, 1021, 1026
 Klein, J.D., xxxi
 Klug, A.C., 24, 129, 337, 516, 1016, 1026
 Knapp, E., 1027
 Knuth, D.E., 303, 438, 1027
 Koch, G., 99, 1027
 Koch, J., 516, 1025
 Kodavalla, H., xxxi
 Kohler, W.H., 1027
 Konopnicki, D., 967, 1027, 1001
 Kornacker, M., 816, 991, 1027
 Korn, F., 888, 1027
 Korth, H.F., 24, 578, 771, 967, 1024, 1026-1027, 1030, 1039, IDOI
 Kossman, D., 771, 888, 1012, 1019, 1027
 Kotidis, Y., 887-888, 1027, 1036
 Koudas, N., 888, 1022
 Koutsoupas, E., 967, 991, 1023-1024
 Kowalski, R.A., 844, 1042
 Kriegel, H-P., 925, 991, 1008-1009, 1011, 1017, 1026, 1037, 1001
 Krishnakumar, N., 771, 1027
 Krishnamurthy, R., 516, 771, 844, 1014, 1016, 1020, 1027
 Krishnaprasad, M., xxxi, 887, 1035
 Kuchenhoff, V., 844, 1042
 Kuhns, J.L., 98, 129, 1027
 Kulkarni, K., xxix
 Kull, D., 691, 1043
 Kumar, K.B., 477, 770, 1016
 Kumar, V., 578, 1027
 Kunchithapadarn, K., xxx
 Kung, H.T., 578, 1027
 Kuo, D., 1027
 Kupsch, J., 1034, 1001
 Kuspert, K., 337, 816, 1028
 LaCroix, M., 129, 1027
 Ladner, R.E., 578, 1030
 Lai, M., 578, 1027
 LaksInnanan, L.V.S., 924, 925, 967, 1027-1028, 1032, 1034
 LaIn, C., 815, 1028
 Lamport, L., 771, 1028
 Lampson, B.W., 771, 1028
 Landers, T.A., 771, 1036
 Landis, G., 815, 1028
 Landwehr, C.L., 722
 Langerak, R., 99, 1028
 Lapis, G., 771, 815, 1022, 1043
 Larson, J.A., 56, 771, 1039
 Larson, P., 390, 438, 887, 1010, 1028, 1035
 Last, M., xxxii
 Lausen, G., 816, 967, 1024, 1026
 Lawande, S., 1029, 1001

- Layman, A., 887, 1022
 Lebowitz, F., 100
 Lee, E.K., 337, 1014
 Lee, M., xxix
 Lee, S., 888, 1044
 Lefebvre, A., 816, 844, 1019, 1042
 Leff, A., 771, 1034
 Lehman, P.L., 578, 1027--1028
 Leinbaugh, P., 1010, 1001
 Lenzerini, M., 56, 1008
 Lescoeur, F., 845, 1036
 Leu, D.F., 1013
 Leung, T.W., 816, 1040
 Leung, T.Y.C., 516, 845, 887, 1028, 1038
 Leventhal, M., 270, 1028
 Levine, F., 578, 602, 1032
 Levy, A.Y., xxxii, 887, 967, 1018-1019, 1040
 Lewis, D., 270, 1028
 Lewis, P.M., 24, 771, 1028, 1036
 Ley, M., xxix
 Libkin, L., 887, 1015
 Liedtke, H., 1042, 1001
 Lieuwen, D.F., 337, 887, 1015, 1025, 1034, 1001
 Lim, E-P., 771, 1028, 1001
 Lin, D., 924, 1028
 Lin, K-I., 991
 Lindsay, B.G., xxxi, 98, 337, 602, 771, 815, 887, 1012, 1022, 1028, 1030-1032, 1041, 1043
 Ling, Y., 516, 1041
 Linnemann, V., 337, 816, 1028
 Lipski, W., 98, 1024
 Lipton, R.J., 1020, 516, 1028, 1008
 Liskov, B., 815, 1029
 Litwin, W., 390, 771, 1029
 Liu, H., 967, 1043
 Liu, M.T., 771, 1017, 1029
 Livny, M., 337, 578, 771, 816, 887, 925, 1006, 1011-1012, 1019, 1023, 1029, 1038, 1044, 1001
 Lochovsky, F., 816, 1026
 Lockemann, P.C., 1042, 1001
 Loo, B., 888, 1007
 Loh, W-Y., 925, 1020
 Lohman, G.M., 516, 691, 771, 815, 1022, 1029, 1042
 Lornet, D.B., 438, 578, 771, 991, 1028-1029, 1033
 Loney, K., 99, 1027
 Loo, B.-L., 771, 1033
 Lorie, R.A., 98, 180, 438, 516, 548, 602, 1007, 1012, 1013, 1017, 1022, 1028-1029, 1038
 Lou, Y., 816, 1029
 Lozinskii, E.L., 845, 1026
 Lucchesi, C.L., 648, 1029
 Lu, H., 770, 1029
 Lu, P., 924, 1043
 Lu, Y., 967, 1012
 Ludaescher, B., 967, 1024
 Lueder, R., 1034, 1001
 Lum, V.Y., 369, 887, 1029, 1040
 Lunt, T., 722, 1029
 Lupash, E., xxxii
 Lyngbaek, P., 815, 1018
 Mackert, L.F., 771, 1029
 MacNicol, R., xxxi
 Madigan, D., 924, 1013
 Mahbod, E., 815, 1018
 Mah, T., 924
 Maheshwari, D., 815, 1029
 Maier, D., 24, 98, 648, 815-816, 844-845, 1008, 1015, 1029-1030, 1044
 Makinouchi, A., 816, 1030
 Manber, D., 578, 1030
 Manku, G., 887, 1030
 Mannila, H., 648, 924-925, 1006, 1014, 1022-1024, 1030, 1041
 Mannino, M.V., 516, 1030
 Manolopoulos, Y., 925, 1018
 Manprempre, C., 966, 1043
 Manthey, H., 99, 1011
 Mark, L., 771, 1029
 Ivarkowitz, V.M., 56, 99, 1030
 Martella, G., 722, 1012
 Maryanski, P., 55, 1034
 Matias, Y., 887-888, 1021
 Matos, V., 129, 816, 1033
 Mattos, N., 181, 816, 1011, 1014
 Ivlaugis, L., 181, 1007
 Ivlauliffe, M.L., 815, 1011
 McCarthy, D.R., 181, 1030
 McCreight, E.M., 369, 1008
 McCune, W.W., 99, 1030
 McGill, M.J., 966, 1037
 Ivlgoveran, D., 99, 1015
 McHugh, J., 967, 1030
 McJones, P.R., 98, 602, 1007, 1022
 McLeod, I., 98, 516, 1006, 1023
 McPherson, J., 337, 815, 1022, 1028
 Ivlecca, G., 816, 967, 1007
 Meenakshi, K., 845, 1008
 Megiddo, N., 887, 1024
 Mehl, J.-N., 98, 180, 1007, 1012-1013
 Ivlehotra, S., 771, 1030
 Mehta, M., 770, 925, 1030, 1038
 Melton, J., xxix, xxxii, 180, 816, 887, 1017, 1030-1031
 Nlenasce, D.A., 771, 1031
 Nlendlzon, A.O., 648, 925, 967, 1007, 1019, 1021, 1029, 1031, 1035, 1001
 Meo, R., 925, 1031
 Meredith, J., 691, 1040
 Merlaldo, P., 967, 1007
 Ivlerlin, P.M., 516, 1013
 Merrett, T.H., 129, 303, 1031
 Meyerson, A., 925, 1033
 Michel, R., 477, 1013
 Michie, D., 925, 1031
 Mihaila, G.A., 967, 1031
 Ivlikkilineni, K.P., 477, 1031
 Ivlliller, R.J., 56, 924, 1031, 1043
 Milne, A.A., 550
 Milo, T., 816, 967, 1009, 1031, 1001
 Minker, J., 98-99, 516, 648, 844, 1007, 1012, 1020, 1031
 Minoura, T., 771, 1031
 Mishra, N., 925, 1022, 1033
 Misra, J., 771, 1013
 Missikoff, M., 691, 1012
 Mitchell, G., 516, 1031
 Moffat, A., 966, 1031, 1043-1044
 Mohan, C., xxix, xxxi, 578, 602, 771, 816, 991, 1027, 1031-1032
 Moran, J., xxxii
 Ivlorimoto, Y., 924, 1019
 Morishita, S., 844, 924, 1016, 1019
 Morris, K.A., 844, 1032
 Morrison, R., 1007
 Motro, A., 56, 1032
 Motwarli, R., 888, 924-925, 1007, 1011, 1013, 1018, 1022, 1033, 1041
 Mukkamala, R., 771, 1032
 Mumick, L.S., 516, 816, 845, 887, 1015, 1022, 1032
 Muntz, R.R., 771, 887, 1028, 1031
 Muralikrishna, M., xxxi, 477, 516, 770, 1016, 1032
 Mutchler, D., 771, 1025
 Muthukrishnan, S., 888
 IvIyers, A.C., 815, 1029

- Myllymaki, J., 1029, 1001
 Nag, B., 1034, 1001
 Naqvi, S.A., 816, 844-845,
 1009, 1011, 1014, 1032
 Narang, I., 578, 1032
 Narasayya, V.R., 691, 888, 1013
 Narayanan, S., 816, 1011
 Nash, O., 338
 Naughton, J.F., 1026, xxix,
 438, 477, 516, 691, 770,
 815-816, 844, 887, 967,
 991, 1005, 1011-1012,
 1016, 1022, 1024, 1028,
 1032, 1034, 1039, 1041,
 1044, 1001
 Navathe, S.B., 24, 55-56, 578,
 924, 1005, 1008, 1017,
 1037, 1039
 Negri, M., 180, 1032
 Neimat, M-A., 390, 815, 1018,
 1029
 Nestorov, S., 925, 967, 1032,
 1041
 Newcomer, E., 548, 1009
 Ng, P., 771, 1043
 Ng, R.T., 337, 888, 924-925,
 1008, 1018, 1025, 1028,
 1032
 Ng, V.T., 925, 1014
 Nguyen, T., 967, 1033, 1037,
 1001
 Nicolas, J-M., 99, 648, 1020
 Nievergelt, J., 390, 991, 1018,
 1033
 Nodine, M.H., 771
 Noga, A., 887
 Nyberg, C., 438, 1033
 Obermarck, R., 771,
 1032, 1033, 1043
 O'Callaghan, L., 925, 1022,
 1033
 Olken, F., 477, 516, 887, 1016,
 1033
 Olshen, R.A., 925, 1010
 Olston, C., 771, 1033, 888, 1007
 Ornićinski, E., 578, 924, 1005,
 1037
 Onassis, A., 889
 O'Neil, E., 24, 1033
 O'Neil, P., 24, 771, 887, 1033
 Ong, K., 816, 1044
 Ooi, B-C., 770, 1029
 Oracle, 651
 Orenstein, J., 815, 1028
 Osborn, S.I., 648, 1029
 Osborne, R., xxxii
 Ozden, B., 1033, 1001
 Ozsoyoglu, G., 129, 516, 722,
 816, 1014, 1024, 1033,
 1001
 Ozsoyoglu, Z.M., 129, 516, 816,
 1029, 1033, 1038
 Ozsu, M.T., 771, 1033
 Page, L., 966, 1011
 Pang, A., 924, 925, 1028, 1032
 Papadimitriou, C.H., 99, 548,
 578, 967, 991, 1023-1024,
 1033
 Papakonstantinou, Y., 771,
 967, 1005, 1033-1034
 Paraboschi, S., 181, 1008, 1012
 Paredaens, J., 648, 1015
 Parent, C., 56, 1039
 Park, J., 516, 887, 1034, 1038
 Patel, J.J., 1034, 1001
 Paton, N., 181, 1016
 Patterson, D.A., 337, 1014,
 1034
 Paul, H., 337, 815, 1034, 1037
 Peckham, J., 55, 1034
 Pei, J., 924, 1023, 1034
 Pelagatti, G., 180, 771, 1012,
 1032
 Petajan, E., 1034, 1001
 Petrov, S.V., 648, 1034
 Petry, F., xxxi
 Pfeffer, A., 816, 991, 1024
 Phipps, G., 844, 1016
 Piatetsky-Shapiro, G., 516,
 924, 1006, 1018, 1034
 Piotr, I., 888
 Pippenger, N., 390, 1018
 Pirahesh, H., 181, 337, 516,
 602, 771, 815, 845, 887,
 1014, 1022, 1028,
 1031-1032, 1034, 1038
 Pirotte, A., 129, 1027
 Pistol, P., 337, 816, 1028
 Pitts-Moultis, N., 270, 1034
 Poosala, V., 516, 888, 1005,
 1008, 1034
 Pope, A., 370
 Popek, G., 1., 98, 1007
 Port, G.S., 845, 1008
 Potarnianos, S., 181, 1040
 Powell, A., 1020
 Pramanik, S., 477, 1019
 Prasad, V.V.V., 924, 1005
 Pregibon, I., 888, 924, 1014,
 1023, 1041
 Drescod, P., 270, 1021
 Price, T.G., 98, 516, 602, 1012,
 1022, 1038
 Prock, A., xxx
 Pruyn, J., xxix
 Psaila, G., 925, 1006, 1031
 Pu, C., 771, 1034
 Putzolu, G.R., 98, 578, 602,
 1007, 1022, 1028
 Qian, X., 887, 1034
 Quass, D., 887, 967, 1030,
 1033-1034
 Quinlan, J.R., 925, 1035
 Rafiei, D., xxxi, 925, 1035
 Raghavan, P., 925, 966, 1006,
 1021
 Raiha, K-J., 648, 1030
 Rajagopalan, S., 887, 1030
 Rajaraman, A., 887, 967, 1023,
 1034
 Ramakrishna, M.V., 390, 1035
 Ramakrishnan, L.V., 845, 1035
 Ramakrishnan, R., 56, 516,
 816, 844-845, 887-888,
 924-925, 991, 1004-1005,
 1008-1010, 1016,
 1021-1021, 1025, 1029,
 1031-1032, 1035, 1038,
 1040, 1044, 1001
 Ramamohanarao, K., 390,
 844-845, 966, 1008, 1035,
 1044
 Ramamritham, K., 548, 578,
 1014, 1024
 Ramamurty, R., xxx
 Raman, B., 888, 1007
 Raman, V., 888, 1007, 1035
 Ramasamy, K., 887, 1016,
 1034, 1039, 1001
 Ramaswamy, S., 888, 1005
 Ranganathan, A., 887, 1035
 Ranganathan, M., 925, 1018
 Ranka, S., 925, 1041
 Rao, P., 845, 1035
 Rae, S.G., 887, 1035
 Rastogi, R., 337, 771, 925,
 1010, 1022, 1025, 1030,
 1033, 1035, 1001
 Rcarnes, I., xxx
 Reed, D.P., 578, 771, 1035
 Reese, G., 219, 1035
 Reeve, C.L., 771., 1009, 1036
 Reina, C., 925, uno
 Reiner, D.S., 516, 1026
 Reisner, P., 180, 1013
 Reiter, R., 98, 1035
 Rengarajan, I., xxxi
 Rescorla, E., 722, 1035
 Reuter, A., 548, 602, 1000,
 1022-1023, 1036
 Richardson, J.E., 815, 337,
 1011-1012
 Rielau, S., 816, 1011
 Rijn, V., 722, 1015
 Riloff, E., 966, 1036, 1001
 Rishe, N., 516, 1041

- Rissanen, J., 648, 925, 1030, 1036
 Rivest, R.L., 390, 722, 1036
 Roberts, G.O., 966
 Robie, J., 967, 1013
 Robinson, J.T., 578, 991, 1019, 1027, 1036
 Rogers, A., 888
 Rohmer, J., 845, 1036
 Rosemall, S., 991, 1018
 Rosenkrantz, D.J., 771, 1036
 Rosenthal, A., 516, 925, 1036-1037, 1041, DOI
 Rosenthal, J.S., 966, 1010
 Ross, K.A., 816, 887-888, 1008, 1013, 1015, 1032, 1036
 Rotem, D., 516, 887, 1033, 1040
 Roth, T., 888, 1007
 Rothnie, J.B., 771, 1009, 1036
 Rousseeuw, P.J., 925, 1025
 Roussopoulos, M., 887, 1036
 Roussopoulos, N., 337, 516, 771, 887, 991, 1013, 1027, 1029, 1036
 Roy, P., 967, 1010
 Rozen, S., 691, 1036
 Rumbaugh, J., 56, 1010, 1036
 Rusinkiewicz, M., 771, 1036, 1001
 Ryan, T.A., 815, 1018
 Sacca, D., 845, 1036
 Sacks-Davis, R., 390, 966, 1035, 1044
 Sadri, F., 967, 1027
 Sagalowicz, D., 1043
 Sager, T., 516, 1030
 Sagiv, Y., 516, 648, 816, 845, 967, 1006, 1008, 1026, 1029, 1034, 1036
 Sagonas K.F., 844-845, 1035, 1037
 Sahuguet, A., 967, 1037
 Salton, G., 966, 1037
 Saluja, S., 924, 1022
 Salveter, S., 516, 1039
 Salzberg, B.J., 303, 337, 438, 578, 991, 1029, 1037
 Sarnarati, P., 722, 1012
 Samet, H., 991, 1037
 Sample, N., 967, 1015
 Sander, J., 925, 1017, 1037
 Sanders, R.E., 219, 1037
 Sandhu, R., 722, 1025
 Saraiya, Y., 844, 1032
 Sarawagi, S., 816, 887, 925, 1005, 1037
 Sathaye, A., xxxii
 Savasere, A., 924, 1037
 Sbatella, L., 180, 1032
 Schauble, P., 966, 1037
 Schek, H.-J., 337, 815, 991, 1034, 1037, 1043
 Schell, R., 722, 1029
 Schiesl, G., xxxii
 Schkolnick, M.M., 98, 578, 691, 1008, 1012, 1018, 1037
 Schlageter, G., 771, 1037
 Schlepphorst, C., 967, 1024
 Schneider, D.A., 390, 438, 477, 516, 770, 1016, 1028-1029
 Schneider, R., 991, 1008, 1011
 Schneier, B., 722, 1037
 Scholl, M.H., 337, 815, 1034, 1037
 Schrefl, M., xxxi
 Schryro, M., 1042, 1001
 Schuh, D.T., 815, 1011
 Schumacher, L., xxix
 Schwarz, P., 602, 771, 1031
 Sciore, E., 516, 648, 1037, 1039, 1001
 Scott, K., 56
 Scott, S., 1019
 Seeger, B., 991, 1008
 Segev, A., 516, 887, 1034, 1038, 1041, 1001
 Seidman, G., 888, 1044
 Selfridge, P.G., 924, 1038
 Selinger, P.G., 98, 180, 516, 602, 722, 771, 1012, 1028, 1038, 1043
 Sellis, T.K., 337, 516, 991, 1018, 1025, 1038
 Seshadri, P., xxix, 516, 816, 844-845, 887, 1014, 1035, 1038, 1040
 Seshadri, S., 477, 516, 1010, 1016, 1022, 1001
 Sevcik, K.C., 888, 991, 1008, 1033
 Shacrlton, M., 967, 1015
 Shafer, J.C., 924925, 1006, 1038
 Shaft, U., xxix xxx, 991, HnO, 1021
 Shah, D., 691, 924, 1012, 1038
 Shamir, A., 722, 1036
 Shan, Ivi-C., 815, 1016, 1018, 1026, 1001
 Shanmuga, mndaram, J., 887, 967, 1012, 1038
 Shapiro, L.D., xxix, 477, 1016, 1038
 Shasha, L., xxix, 578, 691, 771, 1010, 1036, 1038
 Shatkay, H., 925, 1038
 Sheard, T., 99, 1038
 Shekita, E.J., 337, 516, 815, 967, 1011-1012, 1022, 1034, 1038
 Sheldon, M.A., 966, 1043
 Shelloy, P., 924, 1038
 Shenoy, S.T., 516, 1038
 Shepherd, J., 390, 1035
 Sheth, A.P., 56, 771, 1017, 1029, 1036, 1039, 1001, 771
 Shim, K., 816, 887-888, 925, 1013, 1022, 1025, 1035
 Shipman, D.W., 548, 771, 1009, 1036
 Shivakumar, N., 924, 1018
 Shmueli, O., 845, 967, 1009, 1027, 1001
 Shockley, W., 722, 1029
 Shoshani, A., 887, 1038-1039
 Shrira, L., 815, 1029
 Shukla, A., xxix, 887, 1016, 1039, 1044
 Sibley, E.H., 24, 1019
 Siegel, M., 516, 1037, 1039, 1001
 Silberschatz, A., 24, xxx, 337, 578, 771, 10101011, 1025, 1027, 1030, 1033, 1039, 1001
 Silverstein, C., 1011
 Simeon, J., 967, 1010, 1013
 Simoll, A.H., 180, 816, 1031
 Simon, E., 181, 691, 1038-1039
 Simoudis, E., 924, 1018, 1039
 Singhal, A., 771, 1029
 Sistla, A.P., 991, 1024, 1043, 1001
 Skeen, D., 771, 1017, 1039
 Skounakis, M., 1006
 Slack, J.M., xxxi
 Slutz, D.R., 98, 1012
 Smith, D.C.P., 55, 1039
 Smith, J.M., 56, 1039
 Smith, K.P., 337, 722, 1008, 1039
 Slnith, P.D., 303, 1039
 Smyth, P., 924, 1006, 1018, 1030
 Snodgrass, R.T., 181, 816, 844, 1041, 1044, 10CH
 So, B., xxix
 Soda, G., 691, 1012
 Solomon, M.H., 815, 1011
 Soloviev, V., 770, 1030
 Son, S.H., xxxi
 Soparkar, N., 578, 1027, 1039, 1001
 Sorenson, P., 691, 1037
 Spaccapietra, S., 56, 1014, 1039

- Speegle, G., xxxi
 Spencer, L., 925, 1024
 Spertus, E., 966, 1039
 Spiegelhalter, D.J., 925, 1031
 Spiro, P., xxxi
 Spyrtatos, N., 99, 1008
 Srikant, R., 887, 924-925, 1006, 1024, 1039
 Srinivasan, V., 578, 1033, 1039, 1001
 Srivastava, D., 516, 816, 844-845, 887-888, 924, 1035-1036, 1038, 1040
 Srivastava, J., 771, 887, 1028, 1040, 1001
 Stacey, D., 771, 1040
 Stachour, P., 722, 1040
 Stankovic, J.A., 578, 1024, 1040, 1001
 Stavropoulos, H., xxix
 Stearns, R., 771, 1036
 Steel, T.B., 1040
 Stefanescu, M., 967, 1013
 Stemple, D., 99, 1038
 Stewart, M., 438, 1037
 Stokes, L., 516, 1022
 Stolfo, S., 924, 1035
 Stolorz, P., 924, 1006
 Stonebraker, M., 24, 98-99, 181, 337, 477, 691, 771, 815-816, 887-888, 1006, 1016-1017, 1024, 1037, 1040, 1044, 1001
 Stone, C.J., 925, 1010
 Strauss, IvLJ., 888
 Strong, R.R., 390, 1018
 Stuckey, P.I., 516, 845, 1038
 Sturgis, H.E., 771, 1028
 Subrahmanian, V.S., 181, 771, 816, 844, 887, 1005, 1022, 1042, 1044, 1001
 Subralnanian, B., 816, 1040
 Subramanian, I.N., 967, 1027
 Subramanian, S., 967, 1012
 Sueiu, D., xxxii, 967, 1011, 1018, 1031
 Su, J., 816, 1024
 Su, S.I.V., 477, 1031
 Sudarshan, S., 1035, 24, xxix, 337, 516, 816, 844-845, 887, 924, 1025, 1035-1036, 1038-1040, 1001
 Sudkamp, N., 337, 816, 1028
 Sull, W., 516, 1041
 Suri, R., 578, 1041
 Swagerman, R., 816, 1011
 Swarni, A., 516, 924, 1006, 1022, 1041
 Swift, T., 844-845, 1035, 1037, 1041
 Szegedy, M., 887
 Szilagyi, P., 966, 1043
 Tam, B.V., 925, 101A
 Tanaka, H., 477, 1019
 Tanca, L., 181, 844, 1012, 1019
 Tan, C.K., 815, 1011
 Tan, J.S., 887, 1040
 Tan, K-L., 770, 1029
 Tan, V.C., 967, 1018
 Tang, N., xxx
 Tannen, V.B., 816, 1011
 Tansel, A.D., 1041, 1001
 Tatblll, N., 888, 1044
 Tay, I.C., 578, 1041
 Taylor, C.C., 925
 Taylor, C.C., 1031
 Teng, J., xxxi
 Teorey, T.J., 55-56, 99, 1041
 Therber, A., xxx
 Thevenin, J.M., 477, 1013
 Thomas, R.H., 771, 1041
 Thomas, S., 925, 1037, 1041
 Thomas, S.A., 722, 1041
 Thomasian, A., xxxi-xxxii, 568, 578, 1019, 1041
 Thompson, C.R., 771, 1010-1011
 Thuraisingham, B., 722, 1040
 Tiberio, P., 691, 1018
 Tibshirani, R., 924, 1023
 Todd, S.J.P., 98, 1041
 Toivonen, H., 924-925, 1006, 1022, 1030, 1041
 Tokuyama, T., 924, 1019
 Tomasic, A., 966, 1021
 Tompa, F.W., 887, 1010
 Towsley, D., 578, 1024
 rn'aiger, I.L., 98, 548, 602, 771, 1007, 1012, 1017, 1022, 1028, 1041
 Trickey, H., 887, 1015
 Tsangaris, M., 816, 1041
 Tsaparas, P., 966, 1010
 Tsatalos, O.G., 815, 1011
 Tsatsoulis, C., xxxi
 Tsichritsis, D.C., 24, 1026
 Tsotras, V.; xxxii
 Tsou, D., 648, 1041
 Tsukernnan, A., 438, 1037
 rl'sukuda, K., 337, 1008
 Tsur, D., 925, 1041
 Tsur, S., 844-845, 924, 1009, 1011, 1014
 Tucherman, L., 99, 1012
 Tucker, A.B., 24, 1041
 Tufte, K., 1034, 10()}
 Tukcy, J.N., 924, 1041
 'rwi.chell, B.C., 337, 1008
 Ubell, M., xxxi
 Ugur, A., xxxii
 Ullman, J.D., 24, xxx, 56, 98, 303, 390, 516, 648, 844-845, 887, 924-925, 967, 1006, 1008, 1011, 1018, 1020, 1023, 1032, 1034-1035, 1041-1042
 Urban, S.D., 56, 1042
 Uren, S., 438, 1037
 Uthurusamy, R., 924, 1006, 1018
 Valdes, J., 1020, 1001
 Valduriez, P., 691, 771, 1033, 1038
 Valentin, G., 691, 1042
 Van Emden, M., 844, 1042
 Van Gelder, A., 844-845, 1032, 1042
 Van Gucht, D., xxix, 129, 816, 887, 1007, 1035
 Van Rijsbergen, C.J., 966, 1042
 Vance, B., 816, 1011
 Vandenberg, S.I., xxxi, 815-816, 1011, 1040
 Vardi, M.Y., 98, 648, 1021, 1042
 Vaughan, B., 438, 1037
 Vélez, B., 1043
 Vélez, B., 966
 Verkamo, A.I., 924-925, 1006, 1030
 Vianu, V., 24, 98, 648, 8Ui, 844, 967, 1005, 1001
 Vidal, M., 56, 1012
 Vieille, I., 816, 844-845, 1019, 1042
 Viswanathan, S., 1025, 1001
 Vitter, J.S., 888
 Von Bultzingsloewen, G., 516, 1042, 1001
 Von Halle, B., 691, 1019
 Vossen, G., 24, 548, 1042-1043
 Vu, Q., 924, 1039
 \Vade, B.W., 98, 180, 602, 722, 1007, 1012-1013, 1022, 1028
 Wade, N., 966, 1042
 Wagner, R.E., 369, 1042
 Wah, I3.V., 887, 1014
 Walch, G., 337, 816, 1028
 \Valker, A., 771, 845, 1007, 1043
 \VaUrath, M., 337, 816, 1028
 Wang, J., 887, 1014
 Wang, K., 967, 1043
 'Nang, M., xxxii, 888
 \Vang, X.S., 771, 1042

- Wang, H., 888, 102:3
Ward, K., 516, 1021
\Varren, D.S., 844–845, 1030, 1035, 1037, 1041
\Watson, V., 98, 1007
Weber, R., 991, 1043
\Veddell, G.E., 648, 1043
\Vei, J., 1039
\Veihl, W., 602, 1043
Weikum, G., 337, 548, 815, 1034, 1037, 1043
Weiner, J., 967, 1032
\Veinreb, D., 815, 1028
\Veiss, R., 966, 1043
Wenger, K., 1029, 1001
West, IVL, 520
Whitaker, M., xxxii
White, C., 771, 104:3
White, S., 219, 1043
White, S.J., 815, 1011
Widom, J., 24, 99, 181, 771, 887–888, 967, 1006–1007, 1012, 1020–1021, 1030, 1033–1034, 1043–1044
Wiederhold, G., 24, xxix, 303, 337, 771, 887, 1020, 1031, 1034, 1043
Wilkinson, W.K., 438, 477, 578, 1008, 1027
Willett, P., 966, 1025
Williams, R., 771, 1043
Wilms, P.P., 771, 815, 1022, 1043
Wilson, L.O., 924, 1038
\Vinuner, M., 925
\Vimmers, E.L., 925, 1006
\Vinslett, M.S., 99, 722, 1039, 1043
Wiorkowski, G., 691, 1043
\Vise, T.E., 337, 1008
\Vistrand, E., 1006, 1001
\Vitten, L.H., 924, 966, 1043
Woelk, D., 815, 1026
vVolfson, O., 771, 991, 1024, 1043, 1001
\Yong, C.Y., 925, 1014
Wong, E., 516, 771, 1009, 1017, 1036, 1043
Wong, H.K.T., 516, 1020
Wong, L., 816, 1011
'Nang, W., 548, 1009
Wood, D., 477, 1016
Woodruff, A., 1006, 1001
Wright, F.L., 649
Wu, J., 816, 1026
Wylie, K., 888, 1007
Xu, E., 991, 1043
Xu, X., 925, 1017, 1037
Yajima, S., 771, 1025
Yang, D., 56, 99, 1041
Yang, Y., 924, 1043
Yannakakis, Iv!., 516, 1036
Yao, S.B., 578, 1028
Yin, Y., 924, 1023
Yoshikawa, Iv!., 771, 816, 1024–1025
Yossi, M., 887–888
Yost, R.A., 98, 771, 1012, 1043
Young, H.C., 438, 1029
Youssefi, K., 516, 1043
Yuan, L., 816, 1033
Ya, C.T., 771, 1043–1044
Yu, J-B., 991, 1021, 1034, 1001
Yue, K.B., xxxi
Yurttas, S., xxxi
Zaiane, O.R., 924, 1043
Zaki, IvLJ., 924, 1044
Zaniolo, C., 98, 181, 516, 648, 816, 844–845, 1014, 1027, 1036, 1044, 1001
Zaot, M., 925, 1006
Zdonik, S.B., xxix, 516, 816, 888, 925, 1031, 1038, 1040, 1044
Zhang, A., 771, 1017
Zhang, T., 925, 1044
Zhang, W., 1044
Zhao, W., 1040, 1001
Zhao, Y., 887, 1044
Zhou, J., 991, 1043
Zhuge, Y., 887, 1044
Ziauddin, M., xxxi
Zicari, R., 181, 816, 844, 1044, 1001
Zilio, D.C., 691, 1042
Zloof, Iv!M., xxix, 98, 1044
Zobel, J., 966, 1031, 1044
Zukowski, U., 844, 1044
Zuliani, M., 691, 1042
Zwilling, M.J., 815, 1011

SUBJECT INDEX

- 1NF, 615
- 2NF, 619
- 2PC, 759, 761
 - blocking, 760
 - with Presumed Abort, 762
- 2PL, 552
 - distributed databases, 755
- 3NF, 617, 625, 628
- 3PC, 762
- 4NF, 636
- 5NF, 638
- A priori property, 893
- Abandoned privilege, 700
- Abort, 522--523, 533, 535, 583, 593, 759
- Abstract data types, 784--785
- ACA schedule, 530
- Access control, 9, 693-694
- Access invariance, 569
- Access mode in SQL, 538
- Access path, 398
 - most selective, 400
- Access privileges, 695
- Access times for disks, 284, 308
- ACID transactions, 521
- Active databases, 132, 168
- Adding tables in SQL, 91
- Adorned program, 839
- ADTs, 784--785
 - encapsulation, 785
 - storage issues, 799
- Advanced Encryption Standard (AES), 710
- AES, 710
- Aggregate functions in ORDBMSs, 801
- Aggregation in Datalog, 831
- Aggregation in SQL, 151, 164
- Aggregation in the ER model, :39,84
- Algebra
 - relational, 102
- ALTER, 696
- Alternatives for data entries in an index, 276
- Analysis phase of recovery, 580, 588
- ANSI, 6, 58
- API, 195
- Application architectures, 236
- Application programmers, 21
- Application programming interface, 195
- Application servers, 251, 253
- Architecture of a DBMS, 19
- ARIES recovery algorithm, 543, 580, 596
- Armstrong's Axioms, 612
- Array chunks, 800, 870
- Arrays, 781
- Assertions in SQL, 167
- Association rules, 897, 900
 - use for prediction, 902
 - with calendars, 900
 - with item hierarchies, 899
- Asynchronous replication, 741, 750--751, 871
 - Capture and Apply, 752--753
 - change data table (CDT), 753
 - conflict resolution, 751
 - peer-to-peer, 751
 - primary site, 751
- Atomic formulas, 118
- Atomicity, 521-522
- Attribute, 11
- Attribute closure, 614
- Attributes in the ER model, 29
- Attributes in the relational model, 59
- Attributes in XML, 229
- Audit trail, 715
- Authentication, 694
- Authorities, 941
- Authorization, 9, 22
- Authorization graph, 701
- Authorization ID, 697
- Autocommit in JDBC, 198
- AVC set, 909
- AVG, 151
- Avoiding cascading aborts, 530
- Axioms for FDs, 612
- B+ trees, 281, 344
 - bulk-loading, 360
 - deletion, 352
 - for sorting, 4:33
 - height, 345
 - insertion, 348
 - key compression, 358
 - locking, 561
 - order, 345
- search, 347
- selection operation, 442
- sequence set, 345
- B+ trees vs. ISAM, 292
- Bags, 780, 782
- Base table, 87
- BCNF, 616, 622
- Bell-LaPadula security model, 706
- Benchmarks, 506, 683, 691
- Binding
 - early vs. late, 788
- Bioinformatics, 999
- BIRCH, 912
- Birth site, 742
- Bit-sliced signature files, 939
- Bitmap indexes, 866
- Bitmapped join index, 869
- Bitmaps
 - for space management, 317, 328
- Blind writes, 528
- BLOBs, 775, 799
- Block evolution of data, 916
- Block nested loops join, 455
- Blocked I/O, 430
- Blocking, 533, 865
- Blocks in disks, 306
- Bloomjoin, 748
- Boolean queries, 929
- Bounding box, 982
- Boyce-Codd nonnal form, 616, 622
- Buckets, 279
 - in a hashed file, 371
 - in histograms, 486
- Buffer frame, 318
- Buffer management
 - DBMS vs. OS, 322
 - double bufl'ering, 432
 - force approach, 541
 - real systems, 322
 - replacernent policy, 321
 - sequential flooding, 321
 - steal approach, 541
- Buffer manager, 20, 305, 318
 - forcing a page, 323
 - page replacement, 319-320
 - pinning, 319
 - prefetching, 322

- Buffer pool, 318
- Buffered writes, 571
- Building phase in hash join, 46:3
- Bulk data types, 780
- Bulk-loading 13+ trees, 360
- Bushy trees, 415
- Caching of methods, 802
- CADjCA:M, 971
- Calculus
 - relational, 116
- Calendric association rules, 900
- Candidate keys, 29, 64, 76
- Capture and Apply, 752
- Cardinality of a relation, 61
- Cartsian product, 105
- CASCADE in foreign keys, 71
- Cascading aborts, 530
- Cascading operators, 488
- Cascading Style Sheets, 249
- Catalogs, 394–395, 480, 483, 741
- Categorical attribute, 905
- Centralized deadlock detection, 756
- Centralized lock management, 755
- Certification authorities, 712
- CGI, 251
- Chained transactions, 536
- Change data table, 753
- Change detection, 916–917
- Character large object, 776
- Checkpoint, 19, 587
 - fuzzy, 587
- Checkpoints, 543
- Checksum, 307
- Choice of indexes, 653
- Chunking, 800, 870
- Class hierarchies, 37, 83
- Class interface, 806
- Classification, 904–905
- Classification rules, 905
- Classification trees, 906
- Clearance, 706
- Client-server architecture, 237, 738
- CL013, 776
- Clock, 322
- Clock policy, 321
- Close an iterator, 408
- Closure of 1stDs, 612
- CLRs, 584, 592, 596
- Clustered file, 277
- Clustered files, 287
- Clustering, 277, 293, 660, 911
- CODASYL, D.B.T.G., 1014
- Collations in SQL, 140
- Collection hierarchies, 789
- Collection hierarchy, 789
 - for ADT methods, 803
 - real systems, 485
- Collection types, 780
- Collisions, 379
- Column, 59
- Commit, 523, 535, 58:3, 759
- Commit protocols, 751, 758
 - 2PC, 759, 761
 - 3PC, 762
- Communication costs, 739, 744, 749
- Communication protocol, 223
- Compensation log records, 584, 592, 596
- Complete axioms, 613
- Complex types, 779, 795
 - vs. reference types, 795
- Composite search keys, 295, 297
- Compressed histogram, 487
- Compression in B+ trees, 358
- Computer aided design and manufacturing, 971
- Concatenated search keys, 295, 297
- Conceptual design, 13, 27
 - tuning, 669
- Conceptual evaluation strategy, 133
- Conceptual schema, 13
- Concurrency, 9, 17
- Concurrency control
 - multiversion, 572
 - optimistic, 566
 - timestamp, 569
- Concurrent execution, 524
- Conflict equivalence, 550
- Conflict resolution, 751
- Conflict serializability vs. serializability, 561
- Conflict serializable schedule, 550
- Conflicting actions, 526
- Conjunct, 445
 - primary, 399
- Conjunctive normal form (CNF), 398, 445
- Connection pooling, 200
- Connections in .IDBC, 198
- Conservative 2PL, 559
- Consistency, 521
- Content types in XML, 232
- Content-based queries, 972, 988
- Convoy phenomenon, 555
- Cookie, 259
- Cookies, 253
- Coordinator site, 758
- Correlated queries, 147, 504, 506
- Cosine normalization, 932
- Cost estimation, 482–483
- Crash recovery, 9, 18, 22, 541, 580, 583–584, 587–588, 590, 592, 595–596
- Crawler, 939
- CREATE DOMAIN, 166
- CREATE statement
 - SQL, 696
- CREATE TABLE, 62
- CREATE TRIGGER, 169
- CREATE TYPE, 167
- CREATE VIEW, 86
- Creating a relation in SQL, 62
- Critical section, 567
- Cross-product operation, 105
- Cross-tabulation, 855
- C8564 at Wisconsin, xxviii
- CSS, 249
- CUBE operator, 857, 869, 887
- Cursors in SQL, 189, 191
- Cylinders in disks, 306
- Dali, 1001
- Data definition language, 12
- Data Definition Language (DDL), 12, 62, 131
- Data dictionary, 395
- Data Encryption Standard, 710
- Data Encryption Standard (DES), 710
- Data entries in an index, 276
- Data independence, 9, 15, 743
 - distributed, 736
 - logical, 15, 87, 736
 - physical, 15, 736
- Data integration, 995
- Data Manipulation Language, 16
- Data manipulation Language (DML), 131
- Data mining, 7, 849, 889
- Data model, 10
 - multidimensional, 849
 - semantic, 10, 27
- Data partitioning, 730
 - skew, 730
- Data reduction, 747
- Data skew, 730, 733
- Data source, 195
- Data streams, 916
- Data striping in RAID, 309–310
- Data sublanguage, 16
- Data warehouse, 7, 678, 754, 848, 870–871

- dean, 871
- extract, 870
- load, 871
- metadata, 872
- purge, 871
- refresh, 871
- transform, 871
- Database administrators, 21-22
- Database architecture
 - Client-Server vs. Collaborating Servers, 738
- Database design
 - conceptual design, 13, 27
 - for an ORDBMS, 793
 - for OLAP, 853
 - impact of concurrent access, 678
 - normal forms, 615
 - null values, 608
 - physical, 291
 - physical design, 14, 28, 650
 - requirements analysis step, 26
 - role of expected workload, 650
 - role of inclusion dependencies, 639
 - schema refinement, 28, 605
 - tools, 27
 - tuning, 22, 28, 650, 667, 670
- Database management system, 4
- Database tuning, 22, 28, 650, 652, 667
- Databases, 4
- Dataflow for parallelism, 731, 733
- Dataguides, 959
- Datalog, 818-819, 822
 - aggregation, 831
 - comparison with relational algebra, 830
 - input and output, 822
 - least fixpoint, 825-826
 - least model, 824, 826
 - model, 823
 - multiset generation, 832
 - negation, 827-828
 - range-restriction and negation, 828
 - rules, 819
 - safety and range-restriction, 826
 - stratification, 829
- DataSpace, 10CH
- Dates and times in SQL, 140
- DB2
 - Index Advisor, 665
- DBA, 22
- DBI library, 252
- DBMS, 4
 - DBMS architecture, 19
 - DBMS vs. **as**, 322
 - DDL, 12
 - Deadlines
 - hard vs. soft, 994
 - Deadlock, 533
 - detection, 556
 - distributed, 756
 - global vs. local, 756
 - phantom, 757
 - prevention, 558
 - Decision support, 847
 - Decision trees, 906
 - pruning, 907
 - splitting attributes, 907
 - Decompositions, 609
 - dependency-preservation, 621
 - horizontal, 674
 - in the absence of redundancy, 674
 - into 3NF, 625
 - into BCNF, 622
 - lossless-join, 619
 - Decorrelation, 506
 - Decryption, 709
 - Deductions, 820
 - Deductive databases, 820
 - aggregation, 831
 - fixpoint semantics, 824
 - least fixpoint, 826
 - least model, 826
 - least model semantics, 823
 - Magic sets rewriting, 838
 - negation, 827-828
 - optimization, 834
 - repeated inferences, 834
 - Seminaive evaluation, 836
 - unnecessary inferences, 834
 - Deep equality, 790
 - Denormalization, 652, 669, 672
 - Dependency-preserving decomposition, 621
 - Dependent attribute, 904
 - DES, 710
 - Deskstar disk, 308
 - DEVise, 1001
 - Difference operation, 105, 141
 - Digital Libraries project, 997
 - Digital signatures, 713
 - Dimensions, 849
 - Directory
 - of pages, 326
 - of slots, 329
 - Directory doubling, 375
 - Dirty bit, 318
 - Dirty page table, 585, 589
 - Dirty read, 526
 - Discretionary access control, 695
 - Disjunctive selection condition, 445
 - Disk **array**, 309
 - Disk space manager, 21, 304, 316
 - Disk tracks, 306
 - Disks, 305
 - access times, 284, 308
 - blocks, 306
 - controller, 307
 - cylinders, tracks, sectors, 306
 - head, 307
 - physical structure, 306
 - platters, 306
 - Distance function, 911
 - Distinct type in SQL, 167
 - Distributed data independence, 736, 743
 - Distributed databases, 726
 - catalogs, 741
 - commit protocols, 758
 - concurrency control, 755
 - data independence, 743
 - deadlock, 756
 - fragmentation, 739
 - global object names, 742
 - heterogeneous, 737
 - join, 745
 - lock management, 755
 - naming, 741
 - optimization, 749
 - project, 744
 - query processing, 743
 - recovery, 755, 758
 - replication, 741
 - scan, 744
 - select, 744
 - Semijoin and Bloomjoin, 747
 - synchronous vs. asynchronous replication, 750
 - transaction management, 755
 - transparency, 736
 - updates, 750
 - Distributed deadlock, 756
 - Distributed query processing, 743
 - Distributed transaction management, 755
 - Distributed transactions, 736
 - Division, 109
 - in SQL, 150
 - Division operation, 109
 - DML, 16
 - Document type declarations (DTDs), 231
 - Document vector, 930
 - DoD security levels, 708
 - Domain, 29, 59

- Domain constraints, 29, 61, 73, 166
- Domain relational calculus, 122
- Domain-key normal form, 648
- Double buffering, 432
- Drill-down, 854
- Driver, 195–196
 - manager, 195–196
 - types, 196
- DROP, 696
- Dropping tables in SQL, 91
- DTDs, 231
- Duplicates in an index, 278
- Duplicates in SQL, 136
- Durability, 521–522
- Dynamic databases, 560
- Dynamic hashing, 373, 379
- Dynamic indexes, 344, 373, 379
- Dynamic linking, 786
- Dynamic SQL, 194
- Early binding, 788
- Electronic commerce, 221
- Elements in XML, 228
- Embedded SQL, 187
- Encapsulation, 785
- Encryption, 709, 712
- Enforcing integrity constraints, 70
- Entities, 4, 13
- Entity references in XML, 229
- Entity sets in the ER model, 28
- Enumerating alternative plans, 492
- Equality
 - deep vs. shallow, 790
- Equality selection, 292
- Equidepth histogram, 487
- Equijoin, 108
- Equivalence of relational
 - algebra expressions, 414
- Equiwidth histogram, 487
- ER model
 - aggregation, 39, 84
 - attribute domains, 29
 - attributes, 29
 - class hierarchies, 37, 83
 - descriptive attributes, 30
 - entities and entity sets, 28
 - key constraints, 32–33
 - keys, 29
 - overlap and covering, 38
 - participation constraints, 34, 79
- ER rmodel
 - relationships
 - and relationship sets, 29
 - many-to-many, 33
 - many-to-one, 33
 - one-to-many, 33
 - roles, 32
 - weak entities, 35, 82
- ERP, 7**
- Event handler, 247
- Events activating triggers, 168
- Example queries
 - Q1, 110, 120, 123, 137, 145, 147, 154
 - Q2, 112, 120, 123, 139, 146
 - Q3, 113, 139
 - Q4, 113, 139
 - Q5, 113, 141
 - Q6, 114, 142, 149
 - Q7, 115, 121, 123
 - Q8, 115
 - Q9, 116, 121, 124, 150
 - Q10, 116
 - Q11, 117, 123, 135
 - Q12, 119
 - Q13, 120
 - Q14, 121, 124
 - Q15, 134
 - Q16, 138
 - Q17, 140
 - Q18, 140
 - Q19, 143
 - Q20, 144
 - Q21, 146
 - Q22, 148
 - Q23, 148
 - Q24, 149
 - Q25, 151
 - Q26, 151
 - Q27, 152
 - Q28, 152
 - Q29, 153
 - Q30, 153
 - Q31, 154**
 - Q32, 155
 - Q33, 158
 - Q34, 159
 - Q35, 160
 - Q36, 160
 - Q37, 161
- Exclusive locks, 531
- EXEC SQL, 187
- Execution plan, 19
- Expensive predicates, 804
- Exploratory data analysis, 849, 890
- Expressions in SQL, 139, 163
- Expressive power
 - algebra vs. calculus, 124
- Extendible hashing, **373**
 - directory doubling, 375
 - global depth, 376
 - local depth, 377
- Extensibility
 - in an optimizer, 803
 - indexing new types, 800
- Extensible Markup Language (XML), 228, 231–232
- Extensible Style Language (XSL), 228
- External schema, 14
- External sorting, 422, 424, 428, 430, 432, 732
- Failure
 - media, 541, 580
 - system crash, 541, 580
- False positives, 938
- Fan-out, 282, 345, 358–359
- Feature vectors, 970, 972
- Field, 59
- FIFO, 322
- Fifth normal form, 638
- File, 20
 - of records, 275
- File organization, 274
 - clustered, 287
 - hashed, 279
 - indexed, 276
 - random, 284
 - sorted, 285
 - tree, 280
- First in first out (FIFO) policy, 321
- First normal form, 615
- Fixed-length records, 327
- Fixpoint, 824
 - Naive evaluation, 835
 - Seminaive evaluation, 836
- Fixpoint evaluation
 - iterations, 834
- Force vs. no-force, 586
- Force-write, 583, 759
- :Forced reinserts, 985
- Forcing pages, 323, 541, 583
- Foreign key constraints, **66**
- Foreign keys, 76
 - versus aids, 796
- Formulas, 118
- Fourth normal form, 636
- Fragmentation, 739–740
- Frequent itemsets, 893
 - a priori property, 893
- Fully distributed lock
 - management, 756
- Functional dependencies, 611
 - Armstrong's Axioms, 612
 - attribute closure, 614
 - closure, 612
 - minimal cover, 625
 - projecti0I1; 621
- Fuzzy checkpoint, 587
- Gateways, 737
- GenBank, 997
- Generalization, 38

- Generalized Search Trees, 987
- Geographic Information
 - Systems (GIS), 971, 998
- Get next tuple, 408
- GiST, 801, 987
- Global deadlock detection, 756
- Global depth in extendible
 - hashing, 376
- GRANT OPTION, 696
- GRANT statement
 - SQL, 695, 699
- Granting privileges in SQL, 699
- Grid directory, 978
- Grid files, 978
 - convex regions, 981
- Group commit, 996
- Grouping in SQL, 154
- Hash functions, 279, 372, 379, 735
- Hash indexes, 279
- Hash join, 463
 - parallel databases, 733–734
- Hash partitioning, 730
- Hashed files, 279
- Heap files, 20, 276, 284, 324
- Height of a tree, 282, 345
- Heterogeneous databases, 737
 - gateways, 737
- Hierarchical clustering, 912
- Hierarchical data model, 6
- Hierarchical deadlock
 - detection, 757
- Histograms, 485–486
 - compressed, 487
 - equidepth, 487
 - equiwidth, 487
 - real systems, 485
- Horizontal decomposition, 674
- Horizontal fragmentation, 739–740
- Host language, 16, 187
- Hot spots, 535, 674, 678, 680
- HTML, 226, 228, 1001
 - tags, 226
- HTML Fonus, 242
- HTTP
 - absence of state, 258
 - request, 224
 - response, 224
- HTTP protocol, 223
- Hubs, 941
- Hmnan Genome project, 997
- Hybrid hash join, 465
- HyperText Markup Language (HTML), 226, 228
- IBM DB2, 167, 322–323, 327, 331, 333, 357, 359, 422, 446, 452–453, 485, 496, 500, 506, 573, 582, 709, 776, 780, 790, 818, 869, 882
- Iceberg queries, 896
- Identifying owner, 36
- IDS, 6
- Implementation
 - aggregation, 469
 - joins, 455, 457–458, 465
 - hash, 463
 - nested loops, 454
 - projections, 447–449
 - hashing, 449
 - sorting, 448
 - selections, 401, 441–442, 444–446
 - with disjunction, 446
 - B+ tree, 442
 - hash index, 444
 - no disjunction, 445
 - no index, 401, 441–442
 - set-operations, 468
- IMS, 6
- Inclusion dependencies, 639
- Incremental algorithms, 403
- Index, 14, 276
 - duplicate data entries, 278
 - alternatives for data entries, 276
 - B+ tree, 344
 - bitmap, 866
 - clustered vs. unclustered, 277
 - composite key, 295
 - concatenated key, 295
 - data entry, 276
 - dynamic, 344, 373, 379
 - equality query, 295
 - equality vs. range selection, 292
 - extendible hashing, 373
 - fan-out, 282
 - hash, 279, 371
 - buckets, 371
 - hash functions, 372
 - primary and overflow pages, 371
 - in SQL, 299
 - ISAM, 341
 - linear hashing, 379
 - matching a selection, 296, 398
 - multidimensional, 973
 - primary vs. secondary, 277
 - range queries and composite
 - key indexes, 295
 - spatial, 973
 - static, 341
 - static hashing, 371
 - tree, 280
 - unclustered, 288 289
 - unique, 278
- Index advisor, 663
- Index configuration, 663
- Index entries, 339
- Index locking, 561
- Index nested loops join, 402, 457
- Index selection, 653
- Index tuning, 667
- Index-only evaluation, 293, 402
- Index-only plans, 662
- Index-only scan, 452, 471, 495
- Indexes
 - choice, 291
- Indexing new data types, 800
- Inference and security, 715
- Inferences, 820
- Information retrieval, 927
- Informix, 322–323, 327, 331, 333, 359, 422, 446, 452–453, 485, 500, 506, 573, 582, 709, 776, 780, 866, 869
- Informix DDS, 167, 790
- Inheritance hierarchies, 37, 83
- Inheritance in object
 - databases, 787
- Inheritance of attributes, 37
- Insertable-into views, 89
- Instance of a relation, 59
- Instance of a relationship set, 30
- Integration, 995
- Integrity constraints, 9, 12, 32, 34, 38, 63, 79
 - in SQL, 167
 - spatial, 971
 - domain, 61, 73
 - foreign key, 66
 - in SQL, 165–166
 - key, 64
 - transactions in SQL, 72
- Intelligent Miner, 914
- Interface for a class, 806
- Interference, 728
- Internet databases, 7
- Interprocess communication (IPC), 802
- Intersection operation, 104, 141
- Inverse document frequency (IDF), 931
- Inverted indexes, 935
- ISA hierarchies, 37, 899
- ISAM, 292, 341
- ISO, 6, 58
- Isolation, 521
- Isolation level, 199
- Isolation level in SQL, 538
- READ UNCOMMITTED, 539

- REPEATABLE READ, 539
- SERIALIZABLE, 539
- Itemset, 893
 - a priori property, 893
 - frequent, 893
 - support, 893
- Iterations, 834
- Iterator interface, 408
- IVEE, 1001
- Java
 - servlet, 254
- Java Database Connectivity (JDBC), 195, 219, 737, 870
- Java virtual machine, 786
- JavaScript, 245
- JDBC, 195, 198, 219, 737, 870
 - architecture, 196
 - autocommit, 198
 - connection, 198
 - data source, 196
 - Database.IVletaData class, 205
 - driver management, 198
 - driver manager, 195-196
 - Exceptions, 203
 - PreparedStatement class, 200
 - ResultSet class, 201
 - Warnings, 203
- JDBC URL, 198
- JDs, 638
- Join dependencies, 6:38
- Joins, 107
 - Bloomjoin, 748
 - definition, 107
 - distributed databases, 745
 - equijoin, 108
 - implementation, 454, 463
 - block nested loops, 455
 - hybrid hash, 465
 - index nested loops, 457
 - sort-merge, 458
 - natural join, 108
 - outer, 164
 - parallel databases, 732, 734
 - Sernijoin, 747
- KDD, 891
- Key, 29, 611
- Key compression, :358
- Key constraints, 32-33
- Keys
 - candidate, 64, 76
 - candidate vs. search, 280
 - composite search, 295
 - foreign, 76
 - foreign key, **66**
 - primary, 65
- Keys constraints, 64-65
- Knowledge discovery, 890
- Large object, 776
- LastLSN, 585
- Latch, 555
- Late binding, 788
- Least fixpoints, 822, 825
- Least model = least fixpoint, 826
- Least models, 822, 824
- Least recently used (LRU) policy, 321
- Left-deep trees, 415
- Legal relation instance, 63
- Level counter in linear hashing, 379
- Levels of abstraction, 12
- Lexicon, 935
- Linear hashing, 379
 - family of hash functions, 379
 - level counter, 379
- Linear recursion, 831
- Linear scales, 979
- LOB, 776
- Local deadlock detection, 756
- Local depth in extendible hashing, 377
- Locators, 776
- Lock downgrades, 556
- Lock escalation, 566
- Lock manager, 21, 554
 - distributed databases, 755
- Lock upgrade, 555
- Lock-coupling, 562
- Locking, 18
 - downgrading, 556
 - B+ trees, 561
 - concurrency, 678
 - Conservative 2PL, 559
 - distributed databases, 755
 - exclusive locks, 531
 - lock escalation, 566
 - lock upgrade, 555
 - Multiple-granularity, 564
 - performance implications, 678
 - shared locks, 531
 - Strict 2PL, 531
 - update locks, 556
- Locking protocol, 18, 530
- Log, 18, 522, 542, 582
 - abort record, 583
 - commit record, 583
 - compensation record (CLR), 583
 - end record, 583
 - force-write, 583
 - lastLSN, 585
 - pageLSN, 582
 - sequence number (LSN), 582
 - tail, 582
 - update record format, 583
 - WAL, 18
- Log record
 - prevLSN field, 583
 - transID field, 583
 - type field, 583
- Log-based Capture, 752
- Logical data independence, 15, 87, 736
 - views, 15
- Logical schema, 13, 27
- Lossless-join decomposition, 619
- Lost update, 529
- LRU, 322
- Machine learning, 890
- Magic Sets, 506
- Magic sets, 837-838
- Main memory databases, 996
- Nandatory access control, 695
 - objects and subjects, 706
- IVlany-to-many relationship, 33
- Many-to-one relationship, 33
- Market basket, 892
- Markup languages, 226
- Master copy, 751
- Master log record, 587
- Matching index, 398
- Natching phase in hash join, 463
- Materialization of intermediate tables, 407
- Materialization of views, 874
- Materialized views
 - refresh, 876
- MathNIL, 235
- MAX, 151
- .Nlean-time-to-failure, 311
- Ivleasures, 849
- Media failure, 541, 580, 595
- IVledia recovery, 595
- Medical imaging, 971
- Melton
 - ,1.,781
- :Meillory hierarchy, 305
- Merge operator, 731
- Merge sort, 424
- Metadata, 394, 872
- Methods
 - caching, 802
 - interpreted vs. compiled, 802
 - security, 801
- .IVlicrosoft SQL Server, 322-323, 327, 331, 333, 357, 359, 422, 446-447, 452-453, 485, 496, 500, 506, 573, 582, 665, 709, 776, 866, **869**, 882
- MIN, 151
- Mineset, 1001
- Minibase software, 1002

- Minimal cover, 625
- Mirroring in RAID, 313
- Mobile databases, 995
- Model, 823
- Model maintenance, 916
- Modifying a table in SQL, 62
- MOLAP, 850
- Most recently llsed CMRU)
 - policy, 321
- MRP, 7
- .MRU, 322
- Multidatabase system, 737
- Multidimensional data model, 849
- Multilevel relations, 707
- Multilevel transactions, 994
- Multimedia databases, 972, 997
- Multiple-granularity locking, 564
- Multiple-query optimization, 507
- Multisets, 135, 780, 782
- Multivalued dependencies, 634
- Multiversion concurrency control, 572
- MVDs, 634
- Naive fixpoint evaluation, 835
- Named constraints in SQL, 66
- Naming in distributed systems, 741
- Natural join, 108
- Natural language searches, 930
- Nearest neighbor queries, 970
- Negation in Datalog, 828
- Negative border, 919
- Nested collections, 783, 798
- Nested loops join, 454
- Nested queries, 145
 - implementation, 504
- Nested relations
 - nesting, 784
 - unnesting, 783
- Nested transactions, 535, 994
- Nesting operation, 784
- Network data model, 6
- NO AC'rION in foreign keys, 71
- Non-preemptive deadlock prevention, 559
- Nonblocking algorithms, 865
- Nonblocking comrnit protocol, 763
- NonIVolatile storage, :306
- Normal forms, 615
 - 1NF, 615
 - BCNF, 616
 - 2NF, 619
 - :3NF, 617
 - Synthesis, 628
- 4NF, 636
- 5NF, 638
- DKNF, 648
- normalization, 622
- PJNF, 648
- tuning, 669
- Normalization, 622, 652
- Null values, 608
 - implementation, 332
 - in SQL, 67, 69, 71, 162
- Numerical attribute, 905
- Object databases, 12
- Object exchange model (OEM), 947
- Object identifiers, 789
- Object manipulation language, 806
- Object-oriented DBMS, 773, 805, 809
- Object-relational DBMS, 773, 809
- ODBC, 195, 219, 737, 995
- ODL, 805
- ODMG data model
 - attribute, 805
 - class, 805
 - inverse relationship, 805
 - method, 806
 - objects, 805
 - relationship, 805
- OEM, 947
- Oids, 789-790
 - referential integrity, 796
 - versus foreign keys, 796
 - versus URLs, 792
- OLAP, 684, 848-849, 887
 - cross-tabulation, 855
 - database design, 853
 - dimension table, 852
 - fact table, 850
 - pivoting, 855
 - roll-up and drill-down, 854
 - SQL window queries, 859
- OLTP, 847
- OML, 80G
- On-the-fly evaluation, 407
- Olle-to-rnany relationship, 33
- One-to-one relationship, 34
- One-way functions, 710
- Online aggregation, 864
- Online analytic processing (OLAP), 848
- Online transaction processing (OI:rp), 847
- OODBMS vs. ORDBMS, 809
- Opaque types, 785
- Open an iterator, 408
- Open Database Connectivity (ODBC), 195, 219, 737, 995
- Optimistic concurrency control, 566
 - validation, 567
- Optimizers
 - cost estimation, 482
 - real systems, 485
 - decomposing a query into blocks, 479
 - extensibility, 803
 - for OHDBMSs, 803
 - handling expensive predicates, 804
 - histograms, 485
 - nested queries, 504
 - overview, 479
 - real systems, 485, 496, 500, 506
 - relational algebra
 - equivalences, 488
 - rule-based, 507
- OQL, 805, 807
- Oracle, 27, 322-323, :327, 331, 333, 357, 359, 422, 446-447, 452-453, 485, 500, 506, 573, 582, 709, 776, 780, 790, 803, 866, 869, 882
- ORDBMS database design, 793
- ORDBMS implementation, 799
- ORDBMS vs. OODBIVIS, 809
- ORDBMS vs. RDBMS, 809
- Order of a B+ tree, 345
- Outer joins, 164
- Overflow in hash join, 464
- Overlap constraints, 38
- Overloading, 788
- Owner of a weak entity, 36
- Packages in SQL:1999, 131
- Page abstraction, 274, 316
- Page fonnats, 326
 - fixed-length records, 327
 - variable-length records, 328
- Page replacement policy, 318-319, 321
- PageLSN, 582
- Paradise, 1001
- Parallel database architecture
 - shared-memory vs. shared-nothing, 727
- Parallel databases, 726-727
 - blocking, 729
 - bulk loading, 731
 - data partitioning, 729-730
 - interference, 728
 - join, 7:32, 734
 - merge and split, 731
 - optimization, 735
 - pipelining, 729

- scan, 731
- sorting, 732
- speed-up vs. scale-up, 728
- Parameteric query
 - optimization, 507
- Parity, 311
- Partial dependencies, 617
- Partial participation, 34
- Participation constraints, 34, 79
- Partition views, 882
- Partitional clustering, 912
- Partitioned parallelism, 729
- Partitioning, 739
 - hash vs. range, 734
- Partitioning data, 730
- Partitioning phase in hash join, 463
- Path expressions, 781, 948
- Peer-to-peer replication, 751
- Perl modules, 252
- Phantom deadlocks, 757
- Phantom problem, 560, 986
- Phantoms, 538, 559
 - SQL, 538539
- Physical data independence, 15, 736
- Physical database design, 14, 28, 291, 650
- Physical design
 - choices, 652
 - clustered indexes, 293
 - co-clustering, 660
 - index selection, 653
 - index-only plans, 662
 - multiple-attribute indexes, 297
 - nested queries, 677
 - query tuning, 670, 675
 - reducing hot spots, 679
 - role of expected workload, 650
 - tuning queries, 670
 - tuning the choice of indexes, 667
 - tuning the conceptual schema, 669
 - tuning wizard, 663, 665
- Physical schema, 14
- Pin count, 318
- Pinning pages, 319
- Pipelined evaluation, 407, 416, 496
- Pipelined parallelism, 729
- Pivoting, 855
- Platters on disks, 306
- PMML, 891
- Point data, 969
- Pointer swizzling, 802
- Polyinstantiation, 708
- Postings file, 935
- Precedence graph, 551
- Precision, 934
- Precommit, 763
- Predicate locking, 561
- Predictor attribute, 904
 - categorical, 905
 - numerical, 905
- Preemptive deadlock prevention, 559
- Prefetching
 - real systems, 323
- Prefetching pages, 322
- Prepare messages, 759
- Presumed Abort, 762
- PrevLSN, 583
- Primary conjunct in a selection, 399
- Primary copy lock management, 755
- Primary index, 277
- PRIMARY KEY constraint in SQL, 66
- Primary keys, 29, 65
 - in SQL, 66
- Primary page for a bucket, 279
- Primary site replication, 751
- Primary storage, 305
- Primary vs. overflow pages, 371
- Privilege descriptor, 701
- Probing phase in hash join, 463
- Procedural Capture, 753
- Process of knowledge discovery, 891
- Project-join normal form, 648
- Projections, 744
 - definition, 103
 - implementation, 447
- Prolog, 819
- Pruning, 907
- Public-key encryption, 710
- Publish and subscribe, 751
- Pushing selections, 409
- Quantifiers, 118
- Query, 16
- Query block, 479
- Query evaluation plan, 405
- Query language, 16, 73, 100
 - Datalog, 818-819
 - domain relational calculus, 122
 - OQL, 807
 - relational algebra, 102
 - relational completeness, 126
 - SQL, 130
 - tuple relational calculus, 117
 - XQuery, 948
- Query modification, 873
- Query optimization, 404, 507
- bushy trees, 415
- deductive databases, 834
- distributed databases, 749
- enumeration of alternative plans, 492
- left-deep trees, 415
- overview, 405, 479
- parallel databases, 735
- pushing selections, 409
- reduction factors, 483, 485
- relational algebra
 - equivalences, 488
- rule-based, 507
- SQL query block, 479
- statistics, 395
- Query optimizer, 19
- Query pattern, 838
- Query processing
 - distributed databases, 743
- Query tuning, 670
- R trees, 982
 - bounding box, 982
- R+ trees, 986
- RAID, 309-310
 - levels, 310
 - mirroring, 313
 - parity, 311
 - redundancy schemes, 311
 - reliability groups, 312
 - striping unit, 310
- Randomized plan generation, 507
- Range partitioning, 730
- Range queries, 295, 970
- Range selection, 292
- Range-restriction, 826, 828
- Ranked queries, 929
- Raster data, 969
- RDBMS vs. ORDBMS, 809
- Real-time databases, 994
- Recall, 934
- Record formats, 330
 - fixed-length records, 331
 - real systems, 331, 333
 - variable-length records, 331
- Record id, 275, 327
- Record ids
 - real systems, 327
- Records, 11, 60
- Recoverability, 530
- Recoverable schedule, 530, 571
- Recovery, 9, 22, 543, 580
 - Analysis phase, 588
 - ARIES, 580
 - checkpointing, 587
 - compensation log record, 584
 - distributed databases, 755, 758
 - fuzzy checkpoint, 587

- log, 18, 522
- loser transactions, 592
- media failure, 595
- Redo phase, 590
- shadow pages, 596
- three phases of restart, 587
- Undo phase, 592
- update log record, 583
- Recovery manager, 21, 540, 580
- Recursive rules, 818
- Redo phase of recovery, 580, 590
- Reduction factor, 400
- Reduction factors, 483, 485
- Redundancy and anomalies, 607
- Redundancy in RAID, 309
- Redundancy schemes, 311
- Reference types, 795
- Reference types in SQL:1999, 790
- Referential integrity, 70
 - in SQL, 70
 - oids, 796
 - violation options, 70
- Refreshing materialized views, 876
- Region data, 970
- Regression rules, 905
- Regression trees, 906
- Relation, 11, 59
 - cardinality, 61
 - degree, 61
 - instance, 60
 - legal instance, 63
 - schema, 59
- Relational algebra, 103
 - comparison with Datalog, 830
 - division, 109
 - equivalences, 488
 - expression, 102
 - expressive power, 124
 - join, 107
 - projection, 103
 - renaming, 106
 - selection, 103
 - set-operations, 104, 468
- Relational calculus
 - domain, 122
 - expressive power, 124
 - safety, 125
 - tuple, 117
- Relational completeness, 126
- Relational data model, 6
- Relational database
 - instance, 61
 - schema, 61
- Relational model, 10, 57
- Relationships, 4, 13, 29, 33
- Renaming in relational algebra, 106
- Repeating history, 581, 596
- Replacement policy, 318–319
- Replacement sort, 428
- Replication, 739, 741
 - asynchronous, 741, 750–751, 871
 - master copy, 751
 - publish and subscribe, 751
 - synchronous, 741, 750
- Resource managers, 993
- Response time, 524
- Restart after crash, 587
- Result size estimation, 483
- REVOKE statement
 - SQL, 699–700
- Revoking privileges in SQL, 700
- Rid, 275, 327
- Rids
 - real systems, 327
- ROLAP, 852
- Role-based authorization, 697
- Roles in the ER model, 32
- Roll-up, 854
- ROLLUP, 857
- Root of an XML document, 231
- Rotational delay for disks, 308
- Round-robin partitioning, 730
- Row-level triggers, 170
- RSA encryption, 710
- Rule-based query optimization, 507
- Rules in Datalog, 819
- Running information for
 - aggregation, 470
- Runs in sorting, 423
- R* trees, 985
- SABRE, 6
- Safe queries, 125
 - in Datalog, 826
- Safety, 826
- Sampling
 - real systems, 485
- Savepoints, 535
- Scalability, 890
- Scale-up, 728
- Scan, 744
- Schedule, 523
 - avoid cascading abort, 530
 - conflict equivalence, 550
 - conflict serializable, 550
 - recoverable, 530, 571
 - serial, 524
 - serializable, 525, 529
 - strict, 552
 - view serializable, 553
- Schema, 11, 59, 61
- Schema decomposition, 609
- Schema evolution, 669
- Schema refinement, 28, 605
 - denormalization, 672
- Schema tuning, 669
- Search key, 276
- Search space of plans, 492
- Search term, 928
- Second normal form, 619
- Secondary index, 277
- Secondary storage, 305
- Secure Electronic Transaction, 713
- Secure Sockets Layer, 712
- Secure Sockets Layer (SSL), 223
- Security, 22, 694, 696
 - authentication, 694
 - classes, 695, 706
 - discretionary access control, 695
 - encryption, 712
 - inference, 715
 - mandatory access control, 695
 - mechanisms, 693
 - policy, 693
 - privileges, 695
 - statistical databases, 715
 - using views, 704
- Security administrator, 709
- Security levels, 708
- Security of methods, 801
- Seek time for disks, 284, 308
- Selection condition
 - conjunct, 445
 - conjunctive normal form, 445
 - term, 444
- Selection pushing, 409
- Selections, 744
 - definition, 103
- Selectivity
 - of an access path, 399
- Semantic data model, 10, 27
- Semantic integration, 995
- Semijoin, 747
- Semijoin reduction, 747
- Seminaive fixpoint evaluation, 836
- Semistructured data, 946, 1001
- Sequence data, 913
- Sequence of itemsets, 902
- Sequence set in a B+ tree, 345
- Sequential flooding, 321, 472
- Sequential patterns, 901
- Serial schedule, 524
- Serializability, 525, 529, 550, 553, 561
- Serializability graph, 551
- Serializable schedule, 529
- Server-side processing, 254

- Servlet, 254
 - request, 255
 - response, 255
- Servlet interface, 255
- Session key, 712
- Session management, **253**
- Set comparisons in SQL, 148
- SET DEFAULT in foreign keys, 71
- Set operators
 - implementation, 468
 - in relational algebra, 104
 - in SQL, 141
- SET protocol, 713
- Set-difference operation, 105
- SGML, 228
- Shadow page recovery, 596
- Shallow equality, 790
- Shared locks, 531
- Shared-disk architecture, 727
- Shared-memory architecture, 727
- Shared-nothing architecture, 727
- Signature files, 937
- Single-tier architecture, 236
- Skew, 730, 733
- Slot directories, 329
- Snapshots, 753, 882
- Snowflake queries, 869
- SOAP, 222
- Sort-merge join, 403, 458
- Sorted files, 285
- Sorted runs, 423
- Sorting, 732
 - applications, 422
 - blocked I/O, 430
 - double buffering, 432
 - external merge sort
 - algorithm, 424
 - replacement sort, 428
 - using B+ trees, 433
- Sound axioms, 613
- Space-filling curves, 975
- Sparse columns, 866
- Spatial data, 969
 - boundary, 969
 - location, 969
- Spatial extent, 969
- Spatial join queries, 971
- Spatial range queries, 970
- Specialization, 38
- Speed-up, 728
- Split operator, 731
- Split selection, 908
- Splitting attributes, 907
- Splitting vector, 732
- SQL
 - access mode, 538
 - aggregate operations, 164
 - definition, 151
 - implementation, 469
 - ALL, 148, 154
 - ALTER, 696
 - ALTER TABLE, 91
 - ANY, 148, 154
 - AS, 139
 - authorization ID, 697
 - AVG, 151
 - BETWEEN, 657
 - CARDINALITY, 781
 - CASCADE, 71
 - collations, 140
 - COMMIT, 535
 - conformance packages, 131
 - correlated queries, 147
 - COUNT, 151
 - CREATE, 696
 - CREATE DOMAIN, 166
 - CREATE TABLE, 62
 - creating views, 86
 - CUBE, 857
 - cursors, 189
 - holdability, 192
 - ordering rows, 193
 - sensitivity, 192
 - updatability, 191
 - Data Definition Language (DDL), 62, 131
 - Data Manipulation" Language (DML), 131
 - DATE values, 140
 - DELETE, 69
 - DISTINCT, 133, 136
 - DISTINCT for aggregation, 151
 - distinct types, 167
 - DROP, 696
 - DROP TABLE, 91
 - dynamic, 194
 - embedded language
 - programming, 187
 - EXCEPT, 141, 149
 - EXEC, 187
 - EXISTS, 141, 163
 - expressing division, 150
 - expressions, 139, 163
 - giving names to constraints, 66
 - GRANT, 695, 699
 - GRANT OPTION, 696
 - GROUP BY, 154
 - HAVING, 154
 - IN, 141
 - indexing, 299
 - INSERT, 52, 69
 - insertable-into views, 89
- SQL
 - integrity constraints
 - assertions, 69, 167
 - CHECK, 165
 - deferred checking, 72
 - domain constraints, 166
 - effect on modifications, 69
 - PRIMARY KEY, **66**
 - table constraints, 69, 165
 - UNIQUE, 66
 - INTERSECT, 141, 149
 - IS NULL, 163
 - isolation level, 538
 - MAX, 151
 - MIN, 151
 - multisets, 135
- SQL
 - nested subqueries
 - definition, 145
 - implementation, 504
 - NO ACTION, 71
 - NOT, 136
 - null values, 67, 69, 71, 162
 - ORDER BY, 193
 - outer joins, 164
 - phantoms, 538–539
 - privileges, 695
 - DELETE, 696
 - INSERT, **696**
 - REFERENCES, 696
 - SELECT, 695
 - UPDATE, 696
 - query block, 479
 - READ UNCOMMITTED, 539
- SQL
 - referential integrity
 - enforcement, 70
 - REPEATABLE READ, 539
 - REVOKE, 699–700
 - CASCADE, 700
 - ROLLBACK, 535
 - ROLLUP, 857
 - savepoints, 535
 - security, 696
 - SELECT-FROM-WHERE, 133
 - SERIALIZABLE, 539
 - SOME, 148
 - SQLCODE, 191
 - SQLERROR, 189
 - SQLSTATE, 189
 - standardization, 58
 - standards, 180
 - strings, 139
 - SUM, 151
 - transaction support, 535
 - transactions and constraints, 72

- UNION, 141
- UNIQUE, 163
- updatable views, 88
- UPDATE, 63, 69
- view updates, 88
- views, 90
- SQL Server
 - data mining, 914
- SQL/MM
 - Data Mining, 891
 - Framework, 776
 - Full Text, 944
 - Spatial, 969
- SQL/PSM, 212
- SQL/XI\iIL, 948
- SQL:1999, 58, 180, 816, 805
 - array type constructor, 780
 - reference types and oids, 790
 - role-based authorization, 697
 - row type constructor, 780
 - structured types, 780
 - structured user-defined types, 779
 - triggers, 168
- SQL:2003, 180
- SQLCODE, 191
- SQLERROR, 189
- SQLJ, 206
 - iterators, 208
- SQLSTATE, 189
- SRQL, 887
- SSL protocol, 712
- Stable storage, 542, 582
- Standard Generalized Markup Language (SGML), 228
- Standardization, 58
- Star join queries, 869
- Star schema, 853
- Starvation, 554
- Stateless communication
 - protocols, 225
- Statement-level triggers, 170
- Static hashing, 371
- Static indexes, 341
- Statistical databases, 715, 855
- Statistics maintained by DBMS, 395
- Stealing frames, 541
- Stop words, 931
- Storage
 - nonvolatile, 306
 - primary, secondary, and tertiary, 305
 - stable, 542
- Stored procedures, 209
- Storing ADTs and structured types, 799
- Stratification, 829
- comparison to relational algebra, 830
- Streaming data, 916
- Strict 2PL, 530–531, 551, 560
- Strict schedule, 552
- Strings in SQL, 139
- Striping unit, 310
- Structured types, 780
 - storage issues, 799
- Structured user-defined types, 779
- Style sheets, 247
- Subclass, 38
- Substitution principle, 788
- Subtransaction, 755
- SUM, 151
- Superclass, 38
- Superkey, 65, 612
- Support, 893
 - association rule, 897
 - classification and regression, 905
 - frequent itemset, 893
 - itemset sequence, 902
- Swizzling, 802
- Sybase, 27
- Sybase ASE, 322–323, 327, 331, 333, 357, 359, 422, 446–447, 452–453, 485, 500, 506, 573, 582, 709, 776
- Sybase ASIQ, 446, 452–453
- Sybase IQ, 447, 866, 869
- Symmetric encryption, 710
- Synchronous replication, 741, 750
 - read-any write-all technique, 751
 - voting technique, 750
- System catalog, 394
- System catalogs, 12, 330, 395, 480, 483, 741
- System R, 6
- System response time, 524
- System throughput, 524
- Table, 60
- Tags in HTML, 226
- Temporal queries, 999
- Term frequency, 931
- Tertiary storage, 305
- Thin clients, 237
- Third normal form, 617, 625, 628
- Thomas Write Rule, 570
- thrashing, 534
- Three-Phase Commit, 762
- Three-tier architecture, 239
 - middle tier, 240
 - presentation tier, 240
- Throughput, 524
- Time-out for deadlock detection, 757
- Timestamp
 - concurrency control, 569–570
 - buffered writes, 571
 - recoverability, 571
 - deadlock prevention in 2PL, 558
- Tioga, 1001
- Total participation, 34
- TP monitor, 993
- TPC-D, 506
- Tracks in disks, 306
- Trail, 582
- Transaction, 520–521
 - abort, 523
 - blind write, 528
 - commit, 523
 - conflicting actions, 526
 - constraints in SQL, 72
 - customer, 892
 - distributed, 736
 - in SQL, 535
 - locks and performance, 678
 - management in a distributed DBMS, 755
 - multilevel and nested, 994
 - properties, 17, 521
 - read, 523
 - schedule, 523
 - write, 523
- Transaction manager, 21, 541
- Transaction processing
 - monitor, 993
- transaction table, 553, 585, 589
- Transactions
 - nested, 536
 - savepoints, 535
- Transactions and JDBC, 199
- transfer time for disks, 308
- TransID, 583
- Transitive dependencies, 617
- Transparent data distribution, 736
- Travelocity, 6
- Tree-based indexing, 280
- Trees
 - R trees, 982
 - B+ tree, 344
 - classification and regression, 906
 - height, 282
 - ISAM, 341
 - node format for B+ tree, 346
 - Region Quad trees, 976
- Triggers, 132, 168
 - activation, 168
 - row vs. statement level, 170

- use in replication, 753
- Trivial FD, 613
- TSQL, 1001
- Tuning, 28, 650, 652, 667
- tuning for concurrency, 678
- Tuning wizard, 663, 665
- Tuple, 60
- Tuple relational calculus, 117
- Turing award, 6
- Two-Phase Commit, 759, 761
 - Presumed Abort, 762
- Two-phase locking, 552
- Two-tier architecture, 237
- Type constructor, 779
- Type extents, 789
- Types
 - complex vs. reference, 795
 - object equality, 790
- UDDI, 222
- UML, 47
 - class diagrams, 48
 - component diagrams, 49
 - database diagrams, 48
- Undo phase of recovery, 580, 592
- Unicode, 230
- Unified Modeling Language, 47
- Uniform resource identifier (URI), 221
- Union compatibility, 104
- Union operation, 104, 141
- UNIQUE constraint in SQL, 66
- Unique index, 278
- Universal resource locator (CURL), 223
- Unnesting operation, 783
- Unpinning pages, 319
- Unrepeatable read, 528
- Updatable cursors, 191
- Updatable views, 88
- Update locks, 556
- Update log record, 583
- Updates in distributed databases, 750
- Upgrading locks, 555
- URI, 221
- URL, 223
- URLs
 - versus oids, 792
- User-defined aggregates, 801
- User-defined types, 784
- Valid XML documents, 231
- Validation in optimistic **ee**, 567
- Variable-length fields, 332
- Variable-length records, 328
- Vector data, 970
- Vector space model, 930
- Vertical fragmentation, 739---740
- Vertical partitioning, 653
- View maintenance, 876, 881
 - incremental, 877
- View materialization, 874
- View serializability, 553
- View serializable schedule, 553
- Views, 14, 86, 90, 653
 - for security, 704
 - GRANT, 704
 - query modification, 873
 - REVOKE, 704
 - updatable, 88
 - updates on, 88
- VisDB, 1000
- Visualization, 1000
- \Vait-die policy, 558
- Waits-for graph, 556, 756
- WAL, 18, 320, 581, 586
- \Varehouse, 754, 848, 870
- Weak entities, 35, 82
- \Weak entity set, 36
- Web crawler, 939
- Web services, 222
- \Well-formed XML document, 231
- vWindow queries, 859
- Wizard
 - index tuning, 663
- Workflow management, 993
- Workload, 291
- Workloads and database design, 650
- Wound-wait policy, 558
- Write-ahead logging, 18, 320, 581, 586
- WSDL, 222
- XML, 228
 - entity references, 229
 - root, 231
- XML content, 232
- XML DTDs, 231
- XML Schema, 234
- XPath, 250
- XQuery, 948
 - path expressions, 948
- XSL, 228, 250
- XSLT, 250
- Z-order curve, 975

Database Management Systems, known for its practical emphasis and comprehensive coverage, has quickly become one of the leading texts for database courses. The third edition features new material on database application development, with a focus on Internet applications. The hands-on approach introduces students to current standards, including JDBC, XML, and 3-tier application architectures. A new, flexible organization allows instructors to teach either an applications-oriented course or an introductory systems-oriented course. The revised "part" organization with new overview chapters makes it easy to select the chapters you need; in-depth chapters within each part can be optional.

This very current new edition also features pedagogical improvements (e.g., chapter objectives, review questions) and updated and extended discussions of data mining, database tuning wizards, decision support, information retrieval, Internet security, object-oriented databases, transaction processing, and XML data management. The coverage has been revised and expanded throughout to reflect the new SQL:1999 standard.

The new website (www.mhhe.com/ramakrishnan) includes extensive instructor resources:

- Exercise solutions
- Lecture slides in PowerPoint, Postscript, and PDF formats
- Notes on setting up and using Oracle, DB2, and SQL Server
- Online solutions and sample databases for all SQL exercises
- Programming exercises and solutions for database application development
- New online instructor's guide with chapter summaries and teaching tips
- Online chapter on QBE; online end-to-end Case Study

Raghu Ramakrishnan is Professor of Computer Science at the University of Wisconsin, Madison, and a founder and CTO of QUIQ. He is a Fellow of the ACM. Johannes Gehrke is Assistant Professor of Computer Science at Cornell University.



McGraw-Hill EngineeringCS.com

www.McGraw-HillEngineeringCS.com —

Your one-stop online shop for all McGraw-Hill Engineering & Computer Science books, supplemental materials, content, & resources! For the student, the professor, and the professional, this site houses it all for your Engineering and CS needs.

"The third edition of *Database Management Systems* continues the excellent tradition of previous editions. The mixture of theory and practice, along with interesting examples, makes the book engaging and easy to read. SQL is explained effectively—important features are covered without overly complicating the presentation with SQL's arcane details. I am particularly impressed with the well-chosen review questions and exercises, and the extensive bibliography. An excellent addition to any database student's or practitioner's library!" *Jim Melton, Editor of ISO/IEC 9075 (SQL), Oracle Corporation*